

Reviewer Pack — Minimal Order of Tropicalized Brauer Groups and Resolution P...

Entry 29a7a7350450 · INCONCLUSIVE

SEC autonomous research engine — attribution: Ludovico Kubler

2026-06-02 23:09:36 UTC

Minimal Order of Tropicalized Brauer Groups and Resolution Proof Width Correlation

Entry ID: 29a7a7350450 **Verdict:** INCONCLUSIVE **Recorded:** 2026-06-02 23:09:36 UTC

1. Conjecture

Field A (mathematical branch): Tropical Geometry (Tropicalized Brauer Groups) **Field B** (complexity object): Resolution Proof Complexity

Statement:

For any given CNF φ with n variables, the minimal order of its tropicalized Brauer group ($\text{trop}(\text{BrauerGroup}(\varphi))$) is linearly correlated with its resolution proof width $w(\varphi)$, such that $\text{trop}(\text{BrauerGroup}(\varphi)) = \Theta(w(\varphi))$.

Rationale (proposer's reasoning):

Tropicalizing algebraic structures like Brauer groups offers a new perspective on the geometry of these objects. This conjecture posits that the complexity-theoretic measure of resolution proof width can be related to the geometric properties of tropicalized Brauer groups, potentially revealing underlying structural similarities.

Taxonomy category: TropicalGeometry (status at proposal time:)

2. Pre-registration (Popper-style)

Hash (SHA-256 prefix): bf416c35198721af

This hash commits to the conjecture statement, acceptance criterion, and seed list **before** the empirical test runs. Tampering with the test or its data after this hash was computed would be detectable.

Acceptance criterion:

The conjecture is supported if the correlation coefficient between the minimal order of the tropicalized Brauer group and the resolution proof width is at least 0.7 for all seeds, using at least 10 independent random CNFs.

3. Barrier filter (F1)

Final: PASS

Barrier	Verdict	Confidence	LLM ₁	LLM ₂
RELATIVIZATION	UNCERTAIN	1.00	HITS	UNCERTAIN
NATURAL_PROOFS	SAFE	0.50	UNCERTAIN	UNCERTAIN
ALGEBRIZATION	SAFE	0.50	UNCERTAIN	UNCERTAIN
KARP_LIPTON	SAFE	0.50	UNCERTAIN	UNCERTAIN

4. Novelty audit

Verdict: ADJACENT_OK against 9 hits across arXiv + Semantic Scholar + ECCC

Search queries (3): - tropical geometry AND Brauer group AND resolution complexity - resolution proof width AND tropicalized Brauer groups AND CNF - linear correlation AND minimal order AND tropical geometry resolution proof

Top relevant hits considered: - [http://arxiv.org/abs/1207.2443v2] Tropical Teichmuller and Siegel spaces - [http://arxiv.org/abs/1204.3875v2] Tropicalizing vs Compactifying the Torelli morphism - [http://arxiv.org/abs/1204.6154v2] Local Tropicalization - [http://arxiv.org/abs/astro-ph/0208243v7] Measurement of the Flux of Ultrahigh Energy Cosmic Rays from Monocular Observations by the High Resolution Fly's Eye Exp - [http://arxiv.org/abs/astro-ph/0407622v3] A Study of the Composition of Ultra High Energy Cosmic Rays Using the High Resolution Fly's Eye - [http://arxiv.org/abs/1511.06778v1] The 1st Fermi Lat Supernova Remnant Catalog - [http://arxiv.org/abs/1401.1130v3] Event Conditional Correlation: Or How Non-Linear Linear Dependence Can Be - [http://arxiv.org/abs/1003.1544v2] Linear Mappings of Free Algebra

5. Empirical test harness

Multi-seed protocol: 5 seeds (11, 23, 37, 53, 71)

Sandbox: isolated subprocess, stdlib-only, ≤90s wall time per cycle **Execution:** rc=124, elapsed=241.1s

5.1 Generated Python source

```
import random
import math
from fractions import Fraction

def run_trial(seed: int) -> dict:
    random.seed(seed)

    def generate_cnf(n):
        clauses = []
        for _ in range(2**n):
            clause = [random.randint(-n, n) for _ in range(random.randint(1, n))]
            if all(abs(x) != abs(y) for x, y in zip(clause, clause[1:])):
                clauses.append(clause)
        return clauses

    def tropicalize_brauer_group(cnf):
        # Placeholder function to simulate the computation of the Brauer group
        # This is a dummy implementation and does not reflect the actual computation
        return random.randint(1, 10)

    def resolution_proof_width(cnf):
        # Placeholder function to simulate the computation of the resolution proof width
        # This is a dummy implementation and does not reflect the actual computation
        return random.randint(5, 20)
```

```

n_values = [5, 10, 15, 20, 30, 40]
brauer_group_orders = []
proof_widths = []

for n in n_values:
    cnf = generate_cnf(n)
    brauer_group_order = tropicalize_brauer_group(cnf)
    proof_width = resolution_proof_width(cnf)
    brauer_group_orders.append(brauer_group_order)
    proof_widths.append(proof_width)

if len(brauer_group_orders) < 10:
    return {
        "metric_name": "correlation_coefficient",
        "metric_value": None,
        "instances_tested": len(brauer_group_orders),
        "n_max": max(n_values),
        "conjecture_holds": False,
        "counterexample": "not_enough_instances"
    }

correlation_coefficient = sum((x - mean_x) * (y - mean_y) for x, y in zip(brauer_group_orders,
                                math.sqrt(sum((x - mean_x)**2 for x in brauer_group_orders) *
                                                sum((y - mean_y)**2 for y in proof_widths)))
mean_x = sum(brauer_group_orders) / len(brauer_group_orders)
mean_y = sum(proof_widths) / len(proof_widths)

return {
    "metric_name": "correlation_coefficient",
    "metric_value": correlation_coefficient,
    "instances_tested": len(brauer_group_orders),
    "n_max": max(n_values),
    "conjecture_holds": abs(correlation_coefficient) >= 0.7,
    "counterexample": ""
}

if __name__ == "__main__":
    seeds = [int(x) for x in sys.argv[1:]] if len(sys.argv) > 1 else [2, 3, 5, 7, 11, 13, 17, 19,

    results = []
    for seed in seeds:
        result = run_trial(seed)
        print(f"TRIAL: {{\"seed\": {seed}, \"metric_name\": \"{result['metric_name']}\", \"metric_value\": {result['metric_value']}\"}}")
        results.append(result)

    if all("metric_value" in r and r["metric_value"] is not None for r in results):
        mean_value = sum(r["metric_value"] for r in results) / len(results)
        std_value = math.sqrt(sum((r["metric_value"] - mean_value)**2 for r in results) / len(results))
        support_fraction = sum(1 for r in results if abs(r["metric_value"]) >= 0.7) / len(results)
        print(f"RESULT: SUPPORTED mean={mean_value} std={std_value} support_fraction={support_fraction}")
    elif any("counterexample" in r and r["counterexample"] != "" for r in results):
        first_failing_seed = next(r["seed"] for r in results if "counterexample" in r and r["counterexample"] != "")
        print(f"RESULT: FALSIFIED counterexample=\"{next(r['counterexample'] for r in results if 'counterexample' in r and r['counterexample'] != '')}\"")
    else:
        print("RESULT: INCONCLUSIVE insufficient_data")

```

6. Per-seed results

(no seed-level data — test crashed before producing TRIAL: lines)

7. Test stdout (last 2KB)

TIMEOUT after 240s

8. Critic adversarial review

Critic verdict: CHALLENGE

Critic reasoning:

skipped: test crashed before producing data

9. Final verdict & safety rail

Verdict: INCONCLUSIVE

Reasoning:

The test timed out before producing data, which means it did not complete its execution to provide a correlation coefficient. | next: Re-run the test with increased time limits or optimize the code to ensure it completes within the given time frame.

11. Audit log (LLM calls)

Total LLM calls: 10

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
1	propose	ollama_remote	glm4:latest	0	0	15211
2	propose	ollama_remote	glm4:latest	0	0	12626
3	preregistration	ollama_remote	glm4:latest	0	0	10366
4	novelty	ollama_remote	glm4:latest	0	0	8314
5	novelty	ollama_remote	glm4:latest	0	0	10147
6	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	18472
7	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	12542
8	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	26307
9	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	15733
10	judge	ollama_remote	glm4:latest	0	0	27445

Totals: 0 input tokens, 0 output tokens, ~\$0.0000 cost, 157162 ms total latency. Provider mix: {'ollama_remote': 10}

(full prompt+response transcripts available in `research/audit/29a7a7350450.jsonl`)

12. Reproducibility

To reproduce this cycle:

1. Apply the test harness in section 5.1 with seeds 11,23,37,53,71
2. The Python sandbox is pure-stdlib; any Python 3.11+ should suffice
3. The full LLM transcripts are in `research/audit/29a7a7350450.jsonl` for verification of provenance
4. The pre-registration hash in section 2 commits to the test design before execution; tampering would be detectable.

Sandbox file: research/pvsnp_sandbox/test_*.py (per-cycle)

Replay tarball: research/replay/29a7a7350450.tar.gz (if generated)