

Reviewer Pack — Minimal Order of Tropical Monomial Ideals and Resolution Pro...

Entry 38cc27d66bc9 · INCONCLUSIVE

SEC autonomous research engine — attribution: Ludovico Kubler

2026-06-05 08:22:35 UTC

Minimal Order of Tropical Monomial Ideals and Resolution Proof Width

Entry ID: 38cc27d66bc9 **Verdict:** INCONCLUSIVE **Recorded:** 2026-06-05 08:22:35 UTC

1. Conjecture

Field A (mathematical branch): Tropical Geometry **Field B** (complexity object): Resolution Proofs (Proof Complexity)

Statement:

For every Boolean satisfiability instance φ with n variables, the minimal order of the tropical monomial ideal $I(\varphi)$ is linearly correlated with its resolution proof width $w(\varphi)$, such that $\min_order(I(\varphi)) = \Theta(w(\varphi))$.

Rationale (proposer's reasoning):

The study of tropical geometry has shown connections between algebraic varieties and computational problems. Tropical monomial ideals are expected to encode information about the structure of the satisfiability problem, which could be exploited to analyze resolution proof width. A direct correspondence between minimal order and proof width would reveal a novel link between tropical geometry and proof complexity.

Taxonomy category: TROPICAL_GEOMETRY (status at proposal time:)

2. Pre-registration (Popper-style)

Hash (SHA-256 prefix): c5e4846886f2949f

This hash commits to the conjecture statement, acceptance criterion, and seed list **before** the empirical test runs. Tampering with the test or its data after this hash was computed would be detectable.

Acceptance criterion:

The conjecture is supported if, for all Boolean formulas φ with n variables ($n \leq 40$), the ratio between the minimal order of the tropical monomial ideal $I(\varphi)$ and its resolution proof width $w(\varphi)$ falls within the range $[0.5, 2]$ with a correlation coefficient ≥ 0.7 , calculated over at least 30 independent seeds.

3. Barrier filter (F1)

Final: PASS

Barrier	Verdict	Confidence	LLM ₁	LLM ₂
RELATIVIZATION	SAFE	0.50	UNCERTAIN	UNCERTAIN
NATURAL_PROOFS	SAFE	0.50	UNCERTAIN	UNCERTAIN
ALGEBRIZATION	UNCERTAIN	0.80	UNCERTAIN	HITS
KARP_LIPTON	SAFE	0.50	UNCERTAIN	UNCERTAIN

4. Novelty audit

Verdict: NOVEL against 0 hits across arXiv + Semantic Scholar + ECCC

Search queries (3): - 'tropical geometry' AND 'resolution proof width' AND 'minimal order' AND 'monomial ideal' - 'proof complexity' AND 'tropical monomial ideals' AND 'resolution proof width' - min_order('tropical geometry') AND resolution_proofs

5. Empirical test harness

Multi-seed protocol: 5 seeds (11, 23, 37, 53, 71)

Sandbox: isolated subprocess, stdlib-only, ≤90s wall time per cycle **Execution:** rc=124, elapsed=243.7s

5.1 Generated Python source

```
import random
import math

def run_trial(seed: int) -> dict:
    random.seed(seed)

    def generate_random_formula(n):
        clauses = []
        for _ in range(2**n):
            clause = [random.choice([f'x{i+1}', f'~x{i+1}']) for i in range(n)]
            clauses.append(clause)
        return clauses

    def tropical_polynomial(clause):
        return ' + '.join(f'{random.randint(1, 5)}*{var}' for var in clause)

    def tropical_monomial_ideal(clauses):
        ideal = []
        for clause in clauses:
            ideal.append(tropical_polynomial(clause))
        return ' & '.join(ideal)

    def resolution_proof_width(formula):
        # Simplified DPLL-based solver with small-timeout constraints
        timeout = 240
        start_time = time.time()
        if start_time + timeout < time.time():
            return None
        # Placeholder for actual DPLL implementation
        return random.randint(1, 10)

    def minimal_order(tropical_ideal):
        # Simplified calculation of minimal order (placeholder)
```

```

        return len(tropical_ideal.split('&'))

n = random.choice([5, 10, 15, 20, 30, 40])
formula = generate_random_formula(n)
tropical_ideal = tropical_monomial_ideal(formula)
proof_width = resolution_proof_width(formula)
if proof_width is None:
    return {
        "metric_name": "minimal_order",
        "metric_value": -1,
        "instances_tested": 0,
        "n_max": n,
        "conjecture_holds": False,
        "counterexample": "timeout"
    }

min_order = minimal_order(tropical_ideal)
ratio = min_order / proof_width if proof_width != 0 else -1

return {
    "metric_name": "minimal_order",
    "metric_value": ratio,
    "instances_tested": 1,
    "n_max": n,
    "conjecture_holds": 0.5 <= ratio <= 2,
    "counterexample": ""
}

if __name__ == "__main__":
    import time
    import sys

    if not sys.argv[1:]:
        seeds = [2, 3, 5, 7, 11, 13, 17, 19, 23, 29, 31, 37, 41, 43, 47, 53, 59, 61, 67, 71, 73, 79]
    else:
        seeds = [int(s) for s in sys.argv[1:]]

    results = []
    for seed in seeds:
        result = run_trial(seed)
        print(f"TRIAL: {{\"seed\": {seed}, \"metric_name\": \"{result['metric_name']}\", \"metric_value\": {result['metric_value']}\"}}")
        results.append(result)

    if all(r["metric_value"] != -1 for r in results):
        mean = sum(r["metric_value"] for r in results) / len(results)
        std_dev = math.sqrt(sum((r["metric_value"] - mean) ** 2 for r in results) / len(results))
        support_fraction = sum(1 for r in results if r["conjecture_holds"]) / len(results)
        print(f"RESULT: SUPPORTED mean={mean} std={std_dev} support_fraction={support_fraction}")
    else:
        first_failing_seed = next(seed for seed, result in zip(seeds, results) if result["metric_value"] == -1)
        print(f"RESULT: FALSIFIED counterexample=\"timeout\" first_failing_seed={first_failing_seed}")

```

6. Per-seed results

(no seed-level data — test crashed before producing TRIAL: lines)

7. Test stdout (last 2KB)

TIMEOUT after 240s

8. Critic adversarial review

Critic verdict: CHALLENGE

Critic reasoning:

skipped: test crashed before producing data

9. Final verdict & safety rail

Verdict: INCONCLUSIVE

Reasoning:

The test timed out before producing data, which means it did not complete within the expected time frame and thus could not provide evidence to support or falsify the conjecture.
| next: NONE

11. Audit log (LLM calls)

Total LLM calls: 11

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
1	propose	ollama_remote	glm4:latest	0	0	13486
2	propose	ollama_remote	glm4:latest	0	0	12161
3	propose	ollama_remote	glm4:latest	0	0	15535
4	preregistration	ollama_remote	glm4:latest	0	0	13373
5	novelty	ollama_remote	glm4:latest	0	0	21126
6	novelty	ollama_remote	glm4:latest	0	0	8877
7	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	14998
8	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	11282
9	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	11830
10	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	10900
11	judge	ollama_remote	glm4:latest	0	0	66889

Totals: 0 input tokens, 0 output tokens, ~\$0.0000 cost, 200458 ms total latency. Provider mix: {'ollama_remote': 11}

(full prompt+response transcripts available in *research/audit/38cc27d66bc9.jsonl*)

12. Reproducibility

To reproduce this cycle:

1. Apply the test harness in section 5.1 with seeds 11,23,37,53,71
2. The Python sandbox is pure-stdlib; any Python 3.11+ should suffice
3. The full LLM transcripts are in *research/audit/38cc27d66bc9.jsonl* for verification of provenance
4. The pre-registration hash in section 2 commits to the test design before execution; tampering would be detectable.

Sandbox file: *research/pvsnp_sandbox/test_*.py* (per-cycle)

Replay tarball: *research/replay/38cc27d66bc9.tar.gz* (if generated)