

Reviewer Pack — Minimal Order of Quadratic Intersections Bounds Proof Entrop...

Entry 32034ed67bbf · INCONCLUSIVE

SEC autonomous research engine — attribution: Ludovico Kubler

2026-05-30 05:48:30 UTC

Minimal Order of Quadratic Intersections Bounds Proof Entropy for Tseitin Formulas

Entry ID: 32034ed67bbf **Verdict:** INCONCLUSIVE **Recorded:** 2026-05-30 05:48:30 UTC

1. Conjecture

Field A (mathematical branch): Algebraic Geometry (Quadratic Geometry) **Field B** (complexity object): Cumulative Space

Statement:

For a Tseitin formula π on n variables, the minimal order of intersections of two quadratic curves in the projective space over an algebraically closed field, when embedded in the plane, bounds the cumulative proof entropy $\Phi(\pi)$ from above by a multiplicative factor of $2^{O(\log(n))}$.

Rationale (proposer's reasoning):

Quadratic geometry offers a rich structure that could provide new insights into the complexity of Tseitin formulas. The minimal order of quadratic intersections may capture intricate geometric properties that are not visible in simpler algebraic structures, potentially leading to stronger bounds on proof entropy.

Taxonomy category: c003b_cumulative_entropy (status at proposal time:)

2. Pre-registration (Popper-style)

Hash (SHA-256 prefix): ac96e17ba0c51883

This hash commits to the conjecture statement, acceptance criterion, and seed list **before** the empirical test runs. Tampering with the test or its data after this hash was computed would be detectable.

Acceptance criterion:

The conjecture is supported if for all generated Tseitin formulas π on n variables, the minimal order of quadratic intersections bounds the cumulative proof entropy $\Phi(\pi)$ such that $|\Phi(\pi)| \leq 2^{(O(\log(n)))} * \text{min_intersection_order}$.

3. Barrier filter (F1)

Final: PASS

Barrier	Verdict	Confidence	LLM ₁	LLM ₂
RELATIVIZATION	UNCERTAIN	1.00	HITS	UNCERTAIN
NATURAL_PROOFS	SAFE	0.50	UNCERTAIN	UNCERTAIN
ALGEBRIZATION	SAFE	1.00	SAFE	SAFE
KARP_LIPTON	SAFE	1.00	SAFE	UNCERTAIN

4. Novelty audit

Verdict: NOVEL against 0 hits across arXiv + Semantic Scholar + ECCC

5. Empirical test harness

Multi-seed protocol: 5 seeds (11, 23, 37, 53, 71)

Sandbox: isolated subprocess, stdlib-only, ≤90s wall time per cycle **Execution:** rc=0, elapsed=0.4s

5.1 Generated Python source

```
#!/usr/bin/env python3
"""C-003b cumulative entropy – FAITHFUL reference harness.

Computes the REAL cumulative proof entropy
    Phi(pi) = sum_t |activeClauses(sigma_t)|
over an actual resolution refutation (Davis-Putnam variable elimination)
of Tseitin formulas, exactly per the Lean definitions in Conjecture003.lean:
    proofStateEntropy(sigma) := sigma.activeClauses.card
    cumulativeEntropy(states) := sum (map proofStateEntropy states)
plus the width-aware totalLiteralWeight version (sum of clause sizes).

This is NOT the CDCL clause-count proxy used by the old c003b_counterexample.py:
it tracks the active clause database at EVERY elimination step of a genuine
resolution derivation, so it is the actual Phi the formalization names.

Pure stdlib (no networkx/pysat), so it runs in the pipeline sandbox.
Protocol: run_trial(seed)->dict with TRIAL/RESULT lines.

Author: Ludovico Kubler (harness by the SEC engine, hand-verified).
NOTE: no `from __future__ import annotations` – the sandbox injects a prelude
above the file, which would break that statement's must-be-first rule.
"""

import json
import math
import random
import sys

Clause = frozenset # a clause = frozenset of signed ints (var or -var)

# ----- Tseitin formula generation -----

def random_regular_graph(n: int, degree: int, rng: random.Random):
    """Simple random regular graph via configuration model with retries."""
    if (n * degree) % 2 != 0:
        n += 1
    for _ in range(200):
```

```

stubs = []
for v in range(n):
    stubs += [v] * degree
rng.shuffle(stubs)
edges = set()
ok = True
for i in range(0, len(stubs), 2):
    u, w = stubs[i], stubs[i + 1]
    if u == w or (min(u, w), max(u, w)) in edges:
        ok = False
        break
    edges.add((min(u, w), max(u, w)))
if ok and len(edges) == n * degree // 2:
    return n, sorted(edges)
# fallback: cycle + chords
edges = {(i, (i + 1) % n) for i in range(n)}
return n, sorted((min(a, b), max(a, b)) for a, b in edges)

def tseitin_cnf(n: int, edges: list, charge: dict) -> list:
    """Tseitin CNF over edge variables (1-indexed). For each vertex, clauses
    enforce parity of incident edges = charge[v]."""
    evar = {}
    for i, (u, v) in enumerate(edges):
        evar[(u, v)] = i + 1
        evar[(v, u)] = i + 1
    inc = {v: [] for v in range(n)}
    for (u, v) in edges:
        inc[u].append(evar[(u, v)])
        inc[v].append(evar[(u, v)])
    clauses = []
    for v in range(n):
        vars_ = inc[v]
        k = len(vars_)
        tgt = charge.get(v, 0)
        for mask in range(1 << k):
            if bin(mask).count("1") % 2 != tgt:
                cl = []
                for j, var in enumerate(vars_):
                    cl.append(-var if (mask >> j) & 1 else var)
                clauses.append(frozenset(cl))
    return clauses

# ----- Resolution by Davis-Putnam variable elimination -----

def resolve(c1: frozenset, c2: frozenset, var: int):
    """Resolvent of c1,c2 on var (c1 has +var, c2 has -var). None if tautology."""
    r = (c1 - {var}) | (c2 - {-var})
    for lit in r:
        if -lit in r:
            return None # tautology
    return frozenset(r)

def dp_refutation_phi(clauses: list, rng: random.Random):

```

"""Run DP elimination in min-occurrence order; track the active clause DB after each elimination. Returns (phi_count, phi_weight, n_steps, derived_empty)."""

```
phi_count = sum_t |DB_t| (proofStateEntropy = card)
phi_weight = sum_t sum_{C in DB_t} |C| (totalLiteralWeight version)
"""
```

```
db = set(clauses)
variables = set(abs(l) for c in db for l in c)
phi_count = len(db)
phi_weight = sum(len(c) for c in db)
n_steps = 1
derived_empty = frozenset() in db
MAX_DB = 200000 # guard against blow-up on bad instances
```

```
while variables and not derived_empty:
    # min-occurrence elimination order (standard good heuristic)
    occ = {v: 0 for v in variables}
    for c in db:
        for l in c:
            if abs(l) in occ:
                occ[abs(l)] += 1
    var = min(variables, key=lambda v: occ[v])
    variables.discard(var)

    pos = [c for c in db if var in c]
    neg = [c for c in db if -var in c]
    others = [c for c in db if var not in c and -var not in c]

    resolvents = set()
    for cp in pos:
        for cn in neg:
            r = resolve(cp, cn, var)
            if r is not None:
                resolvents.add(r)
                if len(r) == 0:
                    derived_empty = True
    new_db = set(others) | resolvents
    # forward subsumption (keep minimal clauses) – cheap pass
    db = new_db
    phi_count += len(db)
    phi_weight += sum(len(c) for c in db)
    n_steps += 1
    if len(db) > MAX_DB:
        break

return phi_count, phi_weight, n_steps, derived_empty
```

----- Separator (cheap balanced-ish) -----

```
def approx_separator_size(n: int, edges: list) -> int:
    """Cheap proxy for balanced separator size: min vertices whose removal
    splits the graph ~in half. We use a simple BFS-layer cut."""
    adj = {v: set() for v in range(n)}
    for (u, v) in edges:
        adj[u].add(v); adj[v].add(u)
```

```

# BFS from 0, cut at the layer crossing the median
from collections import deque
seen = {0}; layer = {0: 0}; q = deque([0])
while q:
    x = q.popleft()
    for y in adj[x]:
        if y not in seen:
            seen.add(y); layer[y] = layer[x] + 1; q.append(y)
if len(seen) < n: # disconnected
    return 0
half = n // 2
# vertices on the boundary between the first ~half and the rest
order = sorted(range(n), key=lambda v: layer.get(v, 0))
side_a = set(order[:half])
cut = set()
for (u, v) in edges:
    if (u in side_a) != (v in side_a):
        cut.add(u if u not in side_a else v)
return len(cut)

```

----- Trial protocol -----

```

def run_trial(seed: int) -> dict:
    rng = random.Random(seed)
    # sweep several small sizes; report the scaling exponent of Phi vs n
    ns = [6, 8, 10, 12, 14]
    data = []
    for n0 in ns:
        n, edges = random_regular_graph(n0, 3, rng)
        charge = {v: 0 for v in range(n)}
        charge[0] = 1 # odd total -> UNSAT
        clauses = tseitin_cnf(n, edges, charge)
        phi_c, phi_w, steps, empty = dp_refutation_phi(clauses, rng)
        sep = approx_separator_size(n, edges)
        data.append({"n": n, "m_edges": len(edges), "sep": sep,
                    "phi_count": phi_c, "phi_weight": phi_w,
                    "steps": steps, "derived_empty": empty})
    # scaling: log-log slope of phi_count vs n
    xs = [math.log(d["n"]) for d in data if d["phi_count"] > 0]
    ys = [math.log(d["phi_count"]) for d in data if d["phi_count"] > 0]
    if len(xs) >= 2:
        mx = sum(xs) / len(xs); my = sum(ys) / len(ys)
        num = sum((x - mx) * (y - my) for x, y in zip(xs, ys))
        den = sum((x - mx) ** 2 for x in xs)
        slope = num / den if den else 0.0
    else:
        slope = 0.0
    biggest = data[-1]
    # Sub-conjecture under test (SC1): Phi >= sep^2 on the largest instance
    # (a concrete, falsifiable scaling claim derived from the C-003b thesis
    # that cumulative entropy is forced by the separator).
    holds = biggest["phi_count"] >= max(1, biggest["sep"]) ** 2
    return {
        "metric_name": "phi_count_loglog_slope_vs_n",
        "metric_value": round(slope, 4),
    }

```

```

    "instances_tested": len(ns),
    "n_max": max(ns),
    "conjecture_holds": holds,
    "counterexample": "" if holds else
        f"n={biggest['n']} sep={biggest['sep']} phi={biggest['phi_count']} < sep^2",
    "detail": data,
}

if __name__ == "__main__":
    seeds = [int(x) for x in sys.argv[1:]] or [11, 23, 37]
    results = []
    for s in seeds:
        r = run_trial(s)
        results.append(r)
        print("TRIAL: " + json.dumps(r))
    slopes = [r["metric_value"] for r in results]
    holds = sum(1 for r in results if r["conjecture_holds"])
    mean = sum(slopes) / len(slopes)
    if holds == len(results):
        print(f"RESULT: SUPPORTED mean_slope={mean:.4f} support_fraction=1.0")
    elif holds == 0:
        print(f"RESULT: FALSIFIED counterexample=\"phi < sep^2\" first_failing_seed={seeds[0]}")
    else:
        print(f"RESULT: INCONCLUSIVE mixed holds={holds}/{len(results)} mean_slope={mean:.4f}")

```

6. Per-seed results

Seed	Metric value	Holds?	Counterexample
?	2.4348	□	
?	2.0177	□	
?	2.0528	□	
?	2.4751	□	
?	2.4691	□	
?	2.2173	□	
?	2.2983	□	
?	2.7629	□	
?	2.3782	□	
?	2.7571	□	
?	2.6234	□	
?	2.5729	□	
?	2.3297	□	

Aggregate statistics:

Statistic	Value
n_seeds	13
metric_mean	2.4145615384615384
metric_std	0.23518233897149377
metric_ci95_half	0.13045568957619594
metric_min	2.0177
metric_max	2.7629
support_fraction	1.0

7. Test stdout (last 2KB)

```
"m_edges": 15, "sep": 3, "phi_count": 591, "phi_weight": 2826, "steps": 16, "derived_empty": true}, {"  
TRIAL: {"metric_name": "phi_count_loglog_slope_vs_n", "metric_value": 2.2983, "instances_tested": 5,  
TRIAL: {"metric_name": "phi_count_loglog_slope_vs_n", "metric_value": 2.7629, "instances_tested": 5,  
TRIAL: {"metric_name": "phi_count_loglog_slope_vs_n", "metric_value": 2.3782, "instances_tested": 5,
```

8. Critic adversarial review

Critic verdict: CHALLENGE

Critic reasoning:

The test code focuses on the cumulative proof entropy of Tseitin formulas using Davis-Putnam variable elimination, which is not directly related to the conjecture's focus on the minimal order of intersections of quadratic curves. The empirical evidence provided ($n \leq 14$) is too small to establish a strong conclusion about the conjecture's validity.

9. Final verdict & safety rail

Verdict: INCONCLUSIVE

Reasoning:

The test code focuses on cumulative proof entropy using Davis-Putnam variable elimination, which is not directly related to the conjecture's focus on minimal order of intersections of quadratic curves. The empirical evidence provided ($n \leq 14$) is too small to establish a strong conclusion about the conjecture's validity. | next: Further investigation is needed to test the conjecture with larger values of n and to ensure that the metric used in the test is directly related to the conjecture.

11. Audit log (LLM calls)

Total LLM calls: 9

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
1	propose	ollama_remote	glm4:latest	0	0	19568
2	propose	ollama_remote	glm4:latest	0	0	14352
3	propose	ollama_remote	glm4:latest	0	0	16925
4	propose	ollama_remote	glm4:latest	0	0	14102
5	preregistration	ollama_remote	glm4:latest	0	0	12338
6	novelty	ollama_remote	glm4:latest	0	0	9507
7	novelty	ollama_remote	glm4:latest	0	0	9130
8	critic	ollama_remote	glm4:latest	0	0	11402
9	judge	ollama_remote	glm4:latest	0	0	15840

Totals: 0 input tokens, 0 output tokens, ~\$0.0000 cost, 123162 ms total latency. Provider mix: {'ollama_remote': 9}

(full prompt+response transcripts available in `research/audit/32034ed67bbf.jsonl`)

12. Reproducibility

To reproduce this cycle:

1. Apply the test harness in section 5.1 with seeds 11,23,37,53,71
2. The Python sandbox is pure-stdlib; any Python 3.11+ should suffice

3. The full LLM transcripts are in `research/audit/32034ed67bbf.jsonl` for verification of provenance
4. The pre-registration hash in section 2 commits to the test design before execution; tampering would be detectable.

Sandbox file: `research/pvsnp_sandbox/test_*.py` (per-cycle)

Replay tarball: `research/replay/32034ed67bbf.tar.gz` (if generated)