Reviewer Pack — Minimal Volume of Toric Varieties Bounds Resolution Proof Wi…

Entry db71b2b7873c · INCONCLUSIVE

SEC autonomous research engine — attribution: Ludovico Kubler

2026-05-30 05:51:34 UTC

Minimal Volume of Toric Varieties Bounds Resolution Proof Width

Entry ID: `db71b2b7873c` **Verdict**: `INCONCLUSIVE` **Recorded**: 2026-05-30 05:51:34 UTC

1. Conjecture

Field A (mathematical branch): Toric Geometry × Boolean Function Complexity **Field B** (complexity object): Resolution proof complexity

Statement:

> For a given 3-CNF formula F with m clauses and n variables, the minimal volume V(toric) of its associated toric variety in the real projective space is lower-bounded by a linear function of the resolution proof width. Specifically, there exists a constant c > 0 such that for all instances, V(toric)(F) ≥ c * resolution_proof_width(F).

Rationale (proposer's reasoning):

> Toric geometry has been used to represent and analyze certain types of polynomials in complexity theory. The volume of toric varieties is known to be computable from the generating fan and can potentially reflect the complexity of solving a problem, like resolving a proof for a CNF formula. If true, this conjecture would provide a new, geometric perspective on resolution proofs.

Taxonomy category: `c003b_cumulative_entropy` (status at proposal time:)

2. Pre-registration (Popper-style)

Hash (SHA-256 prefix): `04bac54bc78cc4f7`

This hash commits to the conjecture statement, acceptance criterion, and seed list **before** the empirical test runs. Tampering with the test or its data after this hash was computed would be detectable.

Acceptance criterion:

> The conjecture is supported if the mean ratio of V(toric)(F) to c * resolution_proof_width(F) across 30 seeds is ≥ 0.8, with no seed producing a ratio > 1.2.

3. Barrier filter (F1)

Final: `PASS`

Barrier	Verdict	Confidence	LLM₁	LLM₂
RELATIVIZATION	SAFE	0.90	UNCERTAIN	SAFE
NATURAL_PROOFS	SAFE	0.50	UNCERTAIN	UNCERTAIN
ALGEBRIZATION	SAFE	0.50	UNCERTAIN	UNCERTAIN
KARP_LIPTON	SAFE	1.00	SAFE	UNCERTAIN

4. Novelty audit

Verdict: `ADJACENT_OK` against 8 hits across arXiv + Semantic Scholar + ECCC

Search queries (3): - `toric geometry boolean function complexity AND resolution proof complexity` - `resolution proof width lower bound toric variety volume` - `linear function bound on minimal volume toric varieties for resolution proof width`

Top relevant hits considered: - [http://arxiv.org/abs/0911.3482v5] Complexity of Networks (reprise) - [http://arxiv.org/abs/nlin/0101006v1] On Complexity and Emergence - [http://arxiv.org/abs/1407.6169 Multiplicative Complexity of Vector Valued Boolean Functions - [http://arxiv.org/abs/2206.14232v2] The arithmetic volume of hypersurfaces in toric varieties and Mahler measures - [http://arxiv.org/abs/1801.0870 Adaptive Lower Bound for Testing Monotonicity on the Line - [http://arxiv.org/abs/1007.1875v2] Lower Bounds for Quantum Oblivious Transfer - [http://arxiv.org/abs/math/0212044v3] Toric ideals, real toric varieties, and the algebraic moment map - [http://arxiv.org/abs/math/0510615v2] Restriction of A-Discriminants and Dual Defect Toric Varieties

5. Empirical test harness

Multi-seed protocol: 5 seeds (11, 23, 37, 53, 71)
Sandbox: isolated subprocess, stdlib-only, ≤90s wall time per cycle **Execution**: rc=0, elapsed=0.5s

5.1 Generated Python source

```python
#!/usr/bin/env python3
"""C-003b cumulative entropy — FAITHFUL reference harness.

Computes the REAL cumulative proof entropy
    Phi(pi) = sum_t |activeClauses(sigma_t)|
over an actual resolution refutation (Davis-Putnam variable elimination)
of Tseitin formulas, exactly per the Lean definitions in Conjecture003.lean:
    proofStateEntropy(sigma) := sigma.activeClauses.card
    cumulativeEntropy(states) := sum (map proofStateEntropy states)
plus the width-aware totalLiteralWeight version (sum of clause sizes).

This is NOT the CDCL clause-count proxy used by the old c003b_counterexample.py:
it tracks the active clause database at EVERY elimination step of a genuine
resolution derivation, so it is the actual Phi the formalization names.

Pure stdlib (no networkx/pysat), so it runs in the pipeline sandbox.
Protocol: run_trial(seed)->dict with TRIAL/RESULT lines.

Author: Ludovico Kubler (harness by the SEC engine, hand-verified).
NOTE: no `from __future__ import annotations` — the sandbox injects a prelude
above the file, which would break that statement's must-be-first rule.
"""
import json
import math
import random
```

```python
import sys

Clause = frozenset  # a clause = frozenset of signed ints (var or -var)


# ---------- Tseitin formula generation ----------

def random_regular_graph(n: int, degree: int, rng: random.Random):
    """Simple random regular graph via configuration model with retries."""
    if (n * degree) % 2 != 0:
        n += 1
    for _ in range(200):
        stubs = []
        for v in range(n):
            stubs += [v] * degree
        rng.shuffle(stubs)
        edges = set()
        ok = True
        for i in range(0, len(stubs), 2):
            u, w = stubs[i], stubs[i + 1]
            if u == w or (min(u, w), max(u, w)) in edges:
                ok = False
                break
            edges.add((min(u, w), max(u, w)))
        if ok and len(edges) == n * degree // 2:
            return n, sorted(edges)
    # fallback: cycle + chords
    edges = {(i, (i + 1) % n) for i in range(n)}
    return n, sorted((min(a, b), max(a, b)) for a, b in edges)


def tseitin_cnf(n: int, edges: list, charge: dict) -> list:
    """Tseitin CNF over edge variables (1-indexed). For each vertex, clauses
    enforce parity of incident edges = charge[v]."""
    evar = {}
    for i, (u, v) in enumerate(edges):
        evar[(u, v)] = i + 1
        evar[(v, u)] = i + 1
    inc = {v: [] for v in range(n)}
    for (u, v) in edges:
        inc[u].append(evar[(u, v)])
        inc[v].append(evar[(u, v)])
    clauses = []
    for v in range(n):
        vars_ = inc[v]
        k = len(vars_)
        tgt = charge.get(v, 0)
        for mask in range(1 << k):
            if bin(mask).count("1") % 2 != tgt:
                cl = []
                for j, var in enumerate(vars_):
                    cl.append(-var if (mask >> j) & 1 else var)
                clauses.append(frozenset(cl))
    return clauses
```

```python
# ---------- Resolution by Davis-Putnam variable elimination ----------

def resolve(c1: frozenset, c2: frozenset, var: int):
    """Resolvent of c1,c2 on var (c1 has +var, c2 has -var). None if tautology."""
    r = (c1 - {var}) | (c2 - {-var})
    for lit in r:
        if -lit in r:
            return None  # tautology
    return frozenset(r)


def dp_refutation_phi(clauses: list, rng: random.Random):
    """Run DP elimination in min-occurrence order; track the active clause DB
    after each elimination. Returns (phi_count, phi_weight, n_steps, derived_empty).

    phi_count  = sum_t |DB_t|             (proofStateEntropy = card)
    phi_weight = sum_t sum_{C in DB_t}|C| (totalLiteralWeight version)
    """
    db = set(clauses)
    variables = set(abs(l) for c in db for l in c)
    phi_count = len(db)
    phi_weight = sum(len(c) for c in db)
    n_steps = 1
    derived_empty = frozenset() in db
    MAX_DB = 200000  # guard against blow-up on bad instances

    while variables and not derived_empty:
        # min-occurrence elimination order (standard good heuristic)
        occ = {v: 0 for v in variables}
        for c in db:
            for l in c:
                if abs(l) in occ:
                    occ[abs(l)] += 1
        var = min(variables, key=lambda v: occ[v])
        variables.discard(var)

        pos = [c for c in db if var in c]
        neg = [c for c in db if -var in c]
        others = [c for c in db if var not in c and -var not in c]

        resolvents = set()
        for cp in pos:
            for cn in neg:
                r = resolve(cp, cn, var)
                if r is not None:
                    resolvents.add(r)
                    if len(r) == 0:
                        derived_empty = True
        new_db = set(others) | resolvents
        # forward subsumption (keep minimal clauses) — cheap pass
        db = new_db
        phi_count += len(db)
        phi_weight += sum(len(c) for c in db)
        n_steps += 1
        if len(db) > MAX_DB:
            break
```

```
    return phi_count, phi_weight, n_steps, derived_empty


# ---------- Separator (cheap balanced-ish) ----------

def approx_separator_size(n: int, edges: list) -> int:
    """Cheap proxy for balanced separator size: min vertices whose removal
    splits the graph ~in half. We use a simple BFS-layer cut."""
    adj = {v: set() for v in range(n)}
    for (u, v) in edges:
        adj[u].add(v); adj[v].add(u)
    # BFS from 0, cut at the layer crossing the median
    from collections import deque
    seen = {0}; layer = {0: 0}; q = deque([0])
    while q:
        x = q.popleft()
        for y in adj[x]:
            if y not in seen:
                seen.add(y); layer[y] = layer[x] + 1; q.append(y)
    if len(seen) < n:  # disconnected
        return 0
    half = n // 2
    # vertices on the boundary between the first ~half and the rest
    order = sorted(range(n), key=lambda v: layer.get(v, 0))
    side_a = set(order[:half])
    cut = set()
    for (u, v) in edges:
        if (u in side_a) != (v in side_a):
            cut.add(u if u not in side_a else v)
    return len(cut)


# ---------- Trial protocol ----------

def run_trial(seed: int) -> dict:
    rng = random.Random(seed)
    # sweep several small sizes; report the scaling exponent of Phi vs n
    ns = [6, 8, 10, 12, 14]
    data = []
    for n0 in ns:
        n, edges = random_regular_graph(n0, 3, rng)
        charge = {v: 0 for v in range(n)}
        charge[0] = 1  # odd total -> UNSAT
        clauses = tseitin_cnf(n, edges, charge)
        phi_c, phi_w, steps, empty = dp_refutation_phi(clauses, rng)
        sep = approx_separator_size(n, edges)
        data.append({"n": n, "m_edges": len(edges), "sep": sep,
                     "phi_count": phi_c, "phi_weight": phi_w,
                     "steps": steps, "derived_empty": empty})
    # scaling: log-log slope of phi_count vs n
    xs = [math.log(d["n"]) for d in data if d["phi_count"] > 0]
    ys = [math.log(d["phi_count"]) for d in data if d["phi_count"] > 0]
    if len(xs) >= 2:
        mx = sum(xs) / len(xs); my = sum(ys) / len(ys)
        num = sum((x - mx) * (y - my) for x, y in zip(xs, ys))
```

```
        den = sum((x - mx) ** 2 for x in xs)
        slope = num / den if den else 0.0
    else:
        slope = 0.0
    biggest = data[-1]
    # Sub-conjecture under test (SC1): Phi >= sep^2 on the largest instance
    # (a concrete, falsifiable scaling claim derived from the C-003b thesis
    # that cumulative entropy is forced by the separator).
    holds = biggest["phi_count"] >= max(1, biggest["sep"]) ** 2
    return {
        "metric_name": "phi_count_loglog_slope_vs_n",
        "metric_value": round(slope, 4),
        "instances_tested": len(ns),
        "n_max": max(ns),
        "conjecture_holds": holds,
        "counterexample": "" if holds else
            f"n={biggest['n']} sep={biggest['sep']} phi={biggest['phi_count']} < sep^2",
        "detail": data,
    }


if __name__ == "__main__":
    seeds = [int(x) for x in sys.argv[1:]] or [11, 23, 37]
    results = []
    for s in seeds:
        r = run_trial(s)
        results.append(r)
        print("TRIAL: " + json.dumps(r))
    slopes = [r["metric_value"] for r in results]
    holds = sum(1 for r in results if r["conjecture_holds"])
    mean = sum(slopes) / len(slopes)
    if holds == len(results):
        print(f"RESULT: SUPPORTED mean_slope={mean:.4f} support_fraction=1.0")
    elif holds == 0:
        print(f"RESULT: FALSIFIED counterexample=\"phi < sep^2\" first_failing_seed={seeds[0]}")
    else:
        print(f"RESULT: INCONCLUSIVE mixed holds={holds}/{len(results)} mean_slope={mean:.4f}")
```

6. Per-seed results

Seed	Metric value	Holds?	Counterexample
?	2.4348	□	
?	2.0177	□	
?	2.0528	□	
?	2.4751	□	
?	2.4691	□	
?	2.2173	□	
?	2.2983	□	
?	2.7629	□	
?	2.3782	□	
?	2.7571	□	
?	2.6234	□	
?	2.5729	□	
?	2.3297	□	

Aggregate statistics:

Statistic	Value
`n_seeds`	13
`metric_mean`	2.4145615384615384
`metric_std`	0.23518233897149377
`metric_ci95_half`	0.13045568957619594
`metric_min`	2.0177
`metric_max`	2.7629
`support_fraction`	1.0

7. Test stdout (last 2KB)


```
"m_edges": 15, "sep": 3, "phi_count": 591, "phi_weight": 2826, "steps": 16, "derived_empty": true}, {'
TRIAL: {"metric_name": "phi_count_loglog_slope_vs_n", "metric_value": 2.2983, "instances_tested": 5,
TRIAL: {"metric_name": "phi_count_loglog_slope_vs_n", "metric_value": 2.7629, "instances_tested": 5,
TRIAL: {"metric_name": "phi_count_loglog_slope_vs_n", "metric_value": 2.3782, "instances_tested": 5,
```


8. Critic adversarial review

Critic verdict: `CHALLENGE`

Critic reasoning:

> The test code focuses on computing the cumulative proof entropy for resolution refutations of Tseitin formulas, which is not directly related to the minimal volume of toric varieties or resolution proof width. The metric used in the test does not correspond to the mathematical object defined in the conjecture.

9. Final verdict & safety rail

Verdict: `INCONCLUSIVE`

Reasoning:

> The critic challenged the test's relevance to the conjecture, stating that the metric used does not correspond to the mathematical object defined in the conjecture. | next: Further investigation is needed to validate the test's relevance and accuracy before drawing conclusions about the conjecture.

11. Audit log (LLM calls)

Total LLM calls: 6

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
1	propose	ollama_remote	glm4:latest	0	0	19424
2	preregistration	ollama_remote	glm4:latest	0	0	14098
3	novelty	ollama_remote	glm4:latest	0	0	8233
4	novelty	ollama_remote	glm4:latest	0	0	9457
5	critic	ollama_remote	glm4:latest	0	0	11175
6	judge	ollama_remote	glm4:latest	0	0	9048

Totals: 0 input tokens, 0 output tokens, ~$0.0000 cost, 71434 ms total latency. Provider mix: {'ollama_remote': 6}

(full prompt+response transcripts available in `research/audit/db71b2b7873c.jsonl`)

12. Reproducibility

To reproduce this cycle:

1. Apply the test harness in section 5.1 with seeds 11,23,37,53,71
2. The Python sandbox is pure-stdlib; any Python 3.11+ should suffice
3. The full LLM transcripts are in `research/audit/db71b2b7873c.jsonl` for verification of provenance
4. The pre-registration hash in section 2 commits to the test design before execution; tampering would be detectable.

Sandbox file: `research/pvsnp_sandbox/test_*.py` (per-cycle)
Replay tarball: `research/replay/db71b2b7873c.tar.gz` (if generated)