

Reviewer Pack — Minimal Symmetry Group Order in Graphs vs Monotone Circuit L...

Entry 9d8176675b10 · INCONCLUSIVE

SEC autonomous research engine — attribution: Ludovico Kubler

2026-05-22 05:38:09 UTC

Minimal Symmetry Group Order in Graphs vs Monotone Circuit Lower Bounds for k-CLIQUE

Entry ID: 9d8176675b10 **Verdict:** INCONCLUSIVE **Recorded:** 2026-05-22 05:38:09 UTC

1. Conjecture

Field A (mathematical branch): Graph Theory: Symmetry Groups **Field B** (complexity object): Complexity Theory: Monotone Circuit Lower Bounds for k-CLIQUE

Statement:

[‘For any graph G with n vertices, the minimal order of a symmetry group that leaves G invariant is at least $\Theta(n^{\{1/4\}})$.’, ‘This lower bound holds for all graphs on n vertices that are not k-CLIQUEs.’, ‘For k-CLIQUE graphs, the minimal order of a symmetry group is $\Omega(k^2)$.’]

Rationale (proposer’s reasoning):

[‘Graph symmetry has been studied in various contexts but rarely applied to complexity theory. The proposed conjecture links the concept of symmetry in graph theory to circuit complexity by suggesting that highly symmetric graphs require circuits of large size to compute.’, ‘This connection could reveal a new perspective on the design of monotone circuits and potentially lead to new approaches for proving lower bounds.’]

Taxonomy category: MONOTONE_CLIQUE (status at proposal time:)

2. Pre-registration (Popper-style)

Hash (SHA-256 prefix): 220450ea20ea7ee5

This hash commits to the conjecture statement, acceptance criterion, and seed list **before** the empirical test runs. Tampering with the test or its data after this hash was computed would be detectable.

Acceptance criterion:

The conjecture is supported if, among 30 random seeds, at least 80% of non-k-CLIQUE graphs have a minimal symmetry group order within $\Theta(n^{\{1/4\}})$ and no k-CLIQUE graph has a minimal symmetry group order below $\Omega(k^2)$.

3. Barrier filter (F1)

Final: PASS

Barrier	Verdict	Confidence	LLM ₁	LLM ₂
RELATIVIZATION	SAFE	0.50	UNCERTAIN	UNCERTAIN
NATURAL_PROOFS	SAFE	0.50	UNCERTAIN	UNCERTAIN
ALGEBRIZATION	SAFE	0.50	UNCERTAIN	UNCERTAIN
KARP_LIPTON	SAFE	0.50	UNCERTAIN	UNCERTAIN

4. Novelty audit

Verdict: NOVEL against 8 hits across arXiv + Semantic Scholar + ECCC

Search queries (3): - minimal symmetry group order graph theory monotone circuit lower bounds clique-symmetry groups invariant graph theory k-clique lower bounds-graph theory symmetry group lower bound monotone circuits k-clique

Top relevant hits considered: - [<http://arxiv.org/abs/2311.04204v3>] Sharp Thresholds Imply Circuit Lower Bounds: from random 2-SAT to Planted Clique - [<http://arxiv.org/abs/1405.2700v1>] Zero Excess and Minimal Length in Finite Coxeter Groups - [<http://arxiv.org/abs/2504.19966v3>] Quantum circuit lower bounds in the magic hierarchy - [<http://arxiv.org/abs/2306.00214v2>] Topological Symmetry Groups of the Petersen graphs - [<http://arxiv.org/abs/1705.06416v2>] A quasi-isometry invariant and thickness bounds for right-angled Coxeter groups - [<http://arxiv.org/abs/1402.2589v3>] Partitioning Perfect Graphs into Stars - [<http://arxiv.org/abs/1704.06241v3>] On monotone circuits with local oracles and clique lower bounds - [<http://arxiv.org/abs/1406.3065v2>] Lower Bounds for Tropical Circuits and Dynamic Programs

5. Empirical test harness

Multi-seed protocol: 5 seeds (11, 23, 37, 53, 71)

Sandbox: isolated subprocess, stdlib-only, ≤90s wall time per cycle **Execution:** rc=0, elapsed=0.5s

5.1 Generated Python source

```
import random
import math

def is_k_clique(graph, k):
    n = len(graph)
    for i in range(n):
        for j in range(i + 1, n):
            if graph[i][j] == 0:
                return False
    return True

def generate_random_graph(n):
    graph = [[random.randint(0, 1) for _ in range(n)] for _ in range(n)]
    for i in range(n):
        graph[i][i] = 0
    return graph

def is_symmetric(graph):
    n = len(graph)
    for i in range(n):
        for j in range(i + 1, n):
            if graph[i][j] != graph[j][i]:
                return False
    return True
```

```

def find_symmetry_group_order(graph):
    n = len(graph)
    if not is_symmetric(graph):
        return float('inf')

    symmetries = []
    for i in range(n):
        for j in range(i + 1, n):
            if graph[i][j] == 1:
                symmetries.append((i, j))

    order = len(symmetries)
    return order

def run_trial(seed: int) -> dict:
    random.seed(seed)
    n = 40
    instances_tested = 30
    total_order = 0
    k_clique_counterexample = ""

    for _ in range(instances_tested):
        graph = generate_random_graph(n)
        if is_k_clique(graph, n):
            order = find_symmetry_group_order(graph)
            if order < instances_tested:
                k_clique_counterexample = "Graph with small symmetry group found"
                break
        else:
            order = find_symmetry_group_order(graph)
            total_order += order

    mean_order = total_order / instances_tested
    conjecture_holds = (mean_order >= n**0.25) and (k_clique_counterexample == "")

    return {
        "metric_name": "Symmetry Group Order",
        "metric_value": mean_order,
        "instances_tested": instances_tested,
        "conjecture_holds": conjecture_holds,
        "counterexample": k_clique_counterexample
    }

if __name__ == "__main__":
    seeds = [int(x) for x in sys.argv[1:]] or [2, 3, 5, 7, 11, 13, 17, 19, 23, 29, 31, 37, 41, 43, 47]
    results = []

    for seed in seeds:
        result = run_trial(seed)
        print(f"TRIAL: {result}")
        results.append(result)

    mean_order = sum(r["metric_value"] for r in results) / len(results)
    support_fraction = sum(1 for r in results if r["conjecture_holds"]) / len(results)

```

```

if all(r["conjecture_holds"] for r in results):
    print(f"RESULT: SUPPORTED mean={mean_order} std=0.0 support_fraction=1.0")
elif support_fraction >= 0.8:
    print(f"RESULT: SUPPORTED mean={mean_order} std=0.0 support_fraction={support_fraction}")
else:
    for r in results:
        if not r["conjecture_holds"]:
            counterexample = r["counterexample"]
            first_failing_seed = seed
            break

    print(f"RESULT: FALSIFIED counterexample=\"{counterexample}\" first_failing_seed={first_failing_seed}")

```

6. Per-seed results

(no seed-level data — test crashed before producing TRIAL: lines)

7. Test stdout (last 2KB)

```

name': 'Symmetry Group Order', 'metric_value': inf, 'instances_tested': 30, 'conjecture_holds': True,
TRIAL: {'metric_name': 'Symmetry Group Order', 'metric_value': inf, 'instances_tested': 30, 'conjecture_holds': True,
TRIAL: {'metric_name': 'Symmetry Group Order', 'metric_value': inf, 'instances_tested': 30, 'conjecture_holds': True,
TRIAL: {'metric_name': 'Symmetry Group Order', 'metric_value': inf, 'instances_tested': 30, 'conjecture_holds': True,
TRIAL: {'metric_name': 'Symmetry Group Order', 'metric_value': inf, 'instances_tested': 30, 'conjecture_holds': True,
TRIAL: {'metric_name': 'Symmetry Group Order', 'metric_value': inf, 'instances_tested': 30, 'conjecture_holds': True,
TRIAL: {'metric_name': 'Symmetry Group Order', 'metric_value': inf, 'instances_tested': 30, 'conjecture_holds': True,
TRIAL: {'metric_name': 'Symmetry Group Order', 'metric_value': inf, 'instances_tested': 30, 'conjecture_holds': True,
TRIAL: {'metric_name': 'Symmetry Group Order', 'metric_value': inf, 'instances_tested': 30, 'conjecture_holds': True,
TRIAL: {'metric_name': 'Symmetry Group Order', 'metric_value': inf, 'instances_tested': 30, 'conjecture_holds': True,
TRIAL: {'metric_name': 'Symmetry Group Order', 'metric_value': inf, 'instances_tested': 30, 'conjecture_holds': True,
TRIAL: {'metric_name': 'Symmetry Group Order', 'metric_value': inf, 'instances_tested': 30, 'conjecture_holds': True,
TRIAL: {'metric_name': 'Symmetry Group Order', 'metric_value': inf, 'instances_tested': 30, 'conjecture_holds': True,
TRIAL: {'metric_name': 'Symmetry Group Order', 'metric_value': inf, 'instances_tested': 30, 'conjecture_holds': True,
TRIAL: {'metric_name': 'Symmetry Group Order', 'metric_value': inf, 'instances_tested': 30, 'conjecture_holds': True,
TRIAL: {'metric_name': 'Symmetry Group Order', 'metric_value': inf, 'instances_tested': 30, 'conjecture_holds': True,
RESULT: SUPPORTED mean=inf std=0.0 support_fraction=1.0

```

8. Critic adversarial review

Critic verdict: CHALLENGE

Critic reasoning:

The test only involved $n \leq 15$ instances, which is too small to draw a general conclusion about the conjecture's validity. The metric value being infinite for all tested instances does not rule out the possibility that the conjecture fails for larger values of n .

9. Final verdict & safety rail

Verdict: INCONCLUSIVE

Reasoning:

The test only involved $n \leq 15$ instances, which is too small to draw a general conclusion about the conjecture's validity. The metric value being infinite for all tested instances does not rule out the possibility that the conjecture fails for larger values of n . | next: Increase the number of instances tested to at least $n = 30$ and ensure that the test covers a wider range of graph sizes, including k -CLIQUE graphs with varying values of k .

11. Audit log (LLM calls)

Total LLM calls: 11

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
1	propose	ollama_remote	glm4:latest	0	0	14629
2	propose	ollama_remote	glm4:latest	0	0	13549
3	preregistration	ollama_remote	glm4:latest	0	0	9241
4	novelty	ollama_remote	glm4:latest	0	0	8244
5	novelty	ollama_remote	glm4:latest	0	0	9648
6	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	15967
7	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	9798
8	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	10573
9	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	9937
10	critic	ollama_remote	glm4:latest	0	0	14997
11	judge	ollama_remote	glm4:latest	0	0	9396

Totals: 0 input tokens, 0 output tokens, ~\$0.0000 cost, 125978 ms total latency. Provider mix: {'ollama_remote': 11}

(full prompt+response transcripts available in `research/audit/9d8176675b10.jsonl`)

12. Reproducibility

To reproduce this cycle:

1. Apply the test harness in section 5.1 with seeds 11,23,37,53,71
2. The Python sandbox is pure-stdlib; any Python 3.11+ should suffice
3. The full LLM transcripts are in `research/audit/9d8176675b10.jsonl` for verification of provenance
4. The pre-registration hash in section 2 commits to the test design before execution; tampering would be detectable.

Sandbox file: `research/pvsnp_sandbox/test_*.py` (per-cycle)

Replay tarball: `research/replay/9d8176675b10.tar.gz` (if generated)