

Reviewer Pack — Minimal Index of Tropicalized Symplectic Leaves and DPLL Pro...

Entry 9f377a5314c3 · INCONCLUSIVE

SEC autonomous research engine — attribution: Ludovico Kubler

2026-06-07 16:42:26 UTC

Minimal Index of Tropicalized Symplectic Leaves and DPLL Proof Width

Entry ID: 9f377a5314c3 **Verdict:** INCONCLUSIVE **Recorded:** 2026-06-07 16:42:26 UTC

1. Conjecture

Field A (mathematical branch): Tropical Geometry × Symplectic Geometry **Field B** (complexity object): Boolean Satisfiability (DPLL Proof Complexity)

Statement:

For every Boolean satisfiability instance φ with resolution proof width $w(\varphi)$, the minimal index of tropicalized symplectic leaves in its associated symplectic leaf system is linearly correlated with its DPLL proof width, such that the minimal index of tropicalized symplectic leaves = $\Theta(w(\varphi))$.

Rationale (proposer's reasoning):

Symplectic geometry provides a geometric structure on spaces of polynomials that has been previously applied to analyze the complexity of real polynomial systems. Tropicalization offers a bridge between algebraic and geometric structures. If this conjecture holds, it would indicate a deep connection between symplectic geometry and computational complexity by providing a lower bound for DPLL proof width in terms of a geometric invariant.

Taxonomy category: symplectic_tropical_complexity (status at proposal time:)

2. Pre-registration (Popper-style)

Hash (SHA-256 prefix): d1dc237f5e6021f2

This hash commits to the conjecture statement, acceptance criterion, and seed list **before** the empirical test runs. Tampering with the test or its data after this hash was computed would be detectable.

Acceptance criterion:

The minimal index of tropicalized symplectic leaves equals resolution proof width within a factor of 2, measured by the ratio $|\text{minimal_index} - w(\varphi)| \leq 0.5 * w(\varphi)$, across at least 80% of instances with more than 20 variables.

3. Barrier filter (F1)

Final: PASS

Barrier	Verdict	Confidence	LLM ₁	LLM ₂
RELATIVIZATION	UNCERTAIN	1.00	HITS	UNCERTAIN
NATURAL_PROOFS	SAFE	0.50	UNCERTAIN	UNCERTAIN
ALGEBRIZATION	UNCERTAIN	0.90	HITS	UNCERTAIN
KARP_LIPTON	SAFE	0.50	UNCERTAIN	UNCERTAIN

4. Novelty audit

Verdict: ADJACENT_OK against 9 hits across arXiv + Semantic Scholar + ECCC

Search queries (3): - tropical geometry AND symplectic geometry AND DPLL proof complexity - resolution proof width AND tropicalization AND symplectic leaf system - Boolean satisfiability AND minimal index of tropicalized symplectic leaves AND $\theta(w(\phi))$

Top relevant hits considered: - [http://arxiv.org/abs/1612.04764v1] Cohomological aspects on complex and symplectic manifolds - [http://arxiv.org/abs/1206.1925v1] Counting Algebraic Curves with Tropical Geometry - [http://arxiv.org/abs/math/0403015v1] Amoebas of algebraic varieties and tropical geometry - [http://arxiv.org/abs/astro-ph/0208243v7] Measurement of the Flux of Ultra-high Energy Cosmic Rays from Monocular Observations by the High Resolution Fly's Eye Exp - [http://arxiv.org/abs/astro-ph/0407622v3] A Study of the Composition of Ultra High Energy Cosmic Rays Using the High Resolution Fly's Eye - [http://arxiv.org/abs/1511.06778v1] The 1st Fermi Lat Supernova Remnant Catalog - [http://arxiv.org/abs/math/0611768v5] The Invariant Symplectic Action and Decay for Vortices - [http://arxiv.org/abs/2103.16512v2] The tropical symplectic Grassmannian

5. Empirical test harness

Multi-seed protocol: 5 seeds (11, 23, 37, 53, 71)

Sandbox: isolated subprocess, stdlib-only, ≤ 90 s wall time per cycle **Execution:** rc=1, elapsed=0.1s

5.1 Generated Python source

```
import random
import math

def run_trial(seed: int) -> dict:
    random.seed(seed)

    def resolution_width(formula):
        if ' and ' not in formula:
            return 1
        left, right = formula.split(' and ', 1)
        return max(resolution_width(left), resolution_width(right)) + 1

    def generate_formula(n):
        if n == 0:
            return random.choice(['True', 'False'])
        else:
            op = random.choice(['and', 'or'])
            left = generate_formula(n - 1)
            right = generate_formula(n - 1)
            return f'({left} {op} {right})'

    def tropical_index(formula):
        if formula == 'True':
            return 0
```

```

elif formula == 'False':
    return math.inf
else:
    left, op, right = formula[1:-2].split(' ')
    if op == 'and':
        return max(tropical_index(left), tropical_index(right))
    elif op == 'or':
        return min(tropical_index(left), tropical_index(right))

n_max = 0
instances_tested = 0
total_metric_value = 0.0
conjecture_holds = True
counterexample = ""

for n in [5, 10, 15, 20, 30, 40]:
    if n > n_max:
        n_max = n

    for _ in range(5):
        formula = generate_formula(n)
        w_phi = resolution_width(formula)
        t_index = tropical_index(formula)

        instances_tested += 1
        total_metric_value += abs(t_index - w_phi) / w_phi

        if abs(t_index - w_phi) > 0.5 * w_phi:
            conjecture_holds = False
            counterexample = f"Formula: {formula}, T-index: {t_index}, W-phi: {w_phi}"

mean_metric_value = total_metric_value / instances_tested

return {
    "metric_name": "Tropical Index vs Resolution Width Ratio",
    "metric_value": mean_metric_value,
    "instances_tested": instances_tested,
    "n_max": n_max,
    "conjecture_holds": conjecture_holds,
    "counterexample": counterexample
}

if __name__ == "__main__":
    if len(sys.argv) > 1:
        seeds = [int(s) for s in sys.argv[1:]]
    else:
        primes = [2, 3, 5, 7, 11, 13, 17, 19, 23, 29, 31, 37, 41, 43, 47, 53, 59, 61, 67, 71, 73,
        seeds = random.sample(primes, min(30, len(primes)))

    results = []
    for seed in seeds:
        result = run_trial(seed)
        print(f"TRIAL: {result}")
        results.append(result)

    mean_metric_value = sum(r["metric_value"] for r in results) / len(results)

```

```

support_fraction = sum(1 for r in results if r["conjecture_holds"]) / len(results)

if all(r["conjecture_holds"] for r in results):
    print(f"RESULT: SUPPORTED mean={mean_metric_value} std=0.0 support_fraction=1.0")
elif sum(1 for r in results if not r["conjecture_holds"]) / len(results) >= 0.8:
    print(f"RESULT: FALSIFIED counterexample=\"{results[next(i for i, r in enumerate(results)
else:
    print(f"RESULT: INCONCLUSIVE support_fraction={support_fraction}")

```

6. Per-seed results

(no seed-level data — test crashed before producing TRIAL: lines)

7. Test stdout (last 2KB)

--STDERR--

Traceback (most recent call last):

```

File "/home/ludo/Scrivania/SEC/research/pvsnp_sandbox/test_f23fcd2.py", line 75, in <module>
    result = run_trial(seed)
             ^^^^^^^^^^^^^

```

```

File "/home/ludo/Scrivania/SEC/research/pvsnp_sandbox/test_f23fcd2.py", line 45, in run_trial
    w_phi = resolution_width(formula)
             ^^^^^^^^^^^^^^^^^^^^^

```

```

File "/home/ludo/Scrivania/SEC/research/pvsnp_sandbox/test_f23fcd2.py", line 32, in resolution_wi
    left, right = formula.split(' and ')
    ^^^^^^^^^

```

ValueError: too many values to unpack (expected 2)

8. Critic adversarial review

Critic verdict: CHALLENGE

Critic reasoning:

skipped: test crashed before producing data

9. Final verdict & safety rail

Verdict: INCONCLUSIVE

Reasoning:

The test code crashed before producing data, which prevents us from verifying the conjecture's support or falsification. | next: Investigate and fix the crash in the test code to proceed with the verification of the conjecture.

11. Audit log (LLM calls)

Total LLM calls: 9

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
1	propose	ollama_remote	glm4:latest	0	0	13825
2	preregistration	ollama_remote	glm4:latest	0	0	9239
3	novelty	ollama_remote	glm4:latest	0	0	8651
4	novelty	ollama_remote	glm4:latest	0	0	9729
5	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	12450

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
6	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	7959
7	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	9140
8	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	10936
9	judge	ollama_remote	glm4:latest	0	0	12698

Totals: 0 input tokens, 0 output tokens, ~\$0.0000 cost, 94626 ms total latency. Provider mix: {'ollama_remote': 9}

(full prompt+response transcripts available in research/audit/9f377a5314c3.jsonl)

12. Reproducibility

To reproduce this cycle:

1. Apply the test harness in section 5.1 with seeds 11,23,37,53,71
2. The Python sandbox is pure-stdlib; any Python 3.11+ should suffice
3. The full LLM transcripts are in research/audit/9f377a5314c3.jsonl for verification of provenance
4. The pre-registration hash in section 2 commits to the test design before execution; tampering would be detectable.

Sandbox file: research/pvsnp_sandbox/test_*.py (per-cycle)

Replay tarball: research/replay/9f377a5314c3.tar.gz (if generated)