

Reviewer Pack — Minimal Hodge-theoretic Index and Frege Proof Depth Inequality...

Entry 10ae0a5c5991 · INCONCLUSIVE

SEC autonomous research engine — attribution: Ludovico Kubler

2026-06-04 11:12:22 UTC

Minimal Hodge-theoretic Index and Frege Proof Depth Inequality

Entry ID: 10ae0a5c5991 **Verdict:** INCONCLUSIVE **Recorded:** 2026-06-04 11:12:22 UTC

1. Conjecture

Field A (mathematical branch): Algebraic Geometry (Hodge Theory) **Field B** (complexity object): Frege Proof Complexity

Statement:

For every CNF φ with n variables, the minimal Hodge-theoretic index $h(\varphi)$ of its associated algebraic variety X is linearly correlated with its Frege proof depth $d(\varphi)$, such that $h(\varphi) = \Omega(d(\varphi))$. Furthermore, for all CNFs, if $d(\varphi) \geq k$, then $h(\varphi) \leq c \cdot d(\varphi)^2$, where c is a constant.

Rationale (proposer's reasoning):

Hodge theory provides a framework for understanding the geometry of complex algebraic varieties. If the complexity of Frege proofs scales with the geometric properties of the associated variety, this could suggest a deeper connection between geometric and logical complexities.

Taxonomy category: HodgeTheoreticIndexFregeProofDepthInequality (status at proposal time:)

2. Pre-registration (Popper-style)

Hash (SHA-256 prefix): b78aacc6b81cfd0a

This hash commits to the conjecture statement, acceptance criterion, and seed list **before** the empirical test runs. Tampering with the test or its data after this hash was computed would be detectable.

Acceptance criterion:

For every CNF φ with n variables, if the minimal Hodge-theoretic index $h(\varphi)$ is greater than or equal to $\Omega(d(\varphi))$ for all seeds and the ratio of $h(\varphi)$ to $d(\varphi)^2$ squared is less than or equal to a constant c when $d(\varphi) \geq k$, then the conjecture is supported. Falsification occurs if any seed shows $h(\varphi) < \Omega(d(\varphi))$ or $h(\varphi)/d(\varphi)^2 > c$ for some φ with $d(\varphi) \geq k$.

3. Barrier filter (F1)

Final: PASS

Barrier	Verdict	Confidence	LLM ₁	LLM ₂
RELATIVIZATION	SAFE	0.50	UNCERTAIN	UNCERTAIN
NATURAL_PROOFS	SAFE	0.50	UNCERTAIN	UNCERTAIN
ALGEBRIZATION	SAFE	0.50	UNCERTAIN	UNCERTAIN
KARP_LIPTON	SAFE	0.50	UNCERTAIN	UNCERTAIN

4. Novelty audit

Verdict: ADJACENT_OK against 1 hits across arXiv + Semantic Scholar + ECCC

Search queries (3): - "Hodge theory" AND "Frege proof complexity" AND minimal index" - "algebraic variety" [AND] "CNF" [AND] "Frege proof depth" - "Hodge-theoretic index" [AND] "linear correlation" [AND] "Frege proof depth"

Top relevant hits considered: - [s2:f00d943cea727962981a17afc4650e9510d2008f] How a Math Class Can Be Two Places at Once: A Proposed Mathematical/Philosophical Collaboration

5. Empirical test harness

Multi-seed protocol: 5 seeds (11, 23, 37, 53, 71)

Sandbox: isolated subprocess, stdlib-only, ≤90s wall time per cycle **Execution:** rc=0, elapsed=0.1s

5.1 Generated Python source

```
import random
import math
from fractions import Fraction

def run_trial(seed: int) -> dict:
    random.seed(seed)

    def generate_cnf(n):
        clauses = []
        for _ in range(n):
            clause = [random.randint(1, n), -random.randint(1, n)]
            clauses.append(clause)
        return clauses

    def frege_proof_depth(cnf):
        # Simplified Frege proof depth simulation
        return len(cnf) * 2

    def hodge_theoretic_index(cnf):
        # Placeholder for Hodge-theoretic index calculation
        # For simplicity, we use the number of clauses as a proxy
        return len(cnf)

    n_values = [5, 10, 15, 20, 30, 40]
    results = []
    for n in n_values:
        cnf = generate_cnf(n)
        h_index = hodge_theoretic_index(cnf)
        d_depth = frege_proof_depth(cnf)
        results.append({"n": n, "h_index": h_index, "d_depth": d_depth})
```

```

if not results:
    return {
        "metric_name": "Hodge-theoretic Index vs Frege Proof Depth",
        "metric_value": None,
        "instances_tested": 0,
        "n_max": 0,
        "conjecture_holds": False,
        "counterexample": "No instances generated"
    }

h_values = [r["h_index"] for r in results]
d_values = [r["d_depth"] for r in results]

mean_h = sum(h_values) / len(h_values)
mean_d = sum(d_values) / len(d_values)
std_h = math.sqrt(sum((x - mean_h) ** 2 for x in h_values) / len(h_values))
std_d = math.sqrt(sum((x - mean_d) ** 2 for x in d_values) / len(d_values))

correlation = sum((h_values[i] - mean_h) * (d_values[i] - mean_d) for i in range(len(h_values)))

c = 1.0
k = 5

conjecture_holds = all(h >= d for h, d in zip(h_values, d_values)) and all(h <= c * d ** 2 for h, d in zip(h_values, d_values))

return {
    "metric_name": "Hodge-theoretic Index vs Frege Proof Depth",
    "metric_value": correlation,
    "instances_tested": len(results),
    "n_max": max(n_values),
    "conjecture_holds": conjecture_holds,
    "counterexample": ""
}

if __name__ == "__main__":
    import sys
    seeds = [int(s) for s in sys.argv[1:]] or [2, 3, 5, 7, 11, 13, 17, 19, 23, 29, 31, 37, 41, 43, 47]

    results = []
    for seed in seeds:
        result = run_trial(seed)
        print(f"TRIAL: {result}")
        results.append(result)

    mean_metric_value = sum(r["metric_value"] for r in results) / len(results)
    std_metric_value = math.sqrt(sum((r["metric_value"] - mean_metric_value) ** 2 for r in results) / len(results))
    support_fraction = sum(1 for r in results if r["conjecture_holds"]) / len(results)

    if all(r["conjecture_holds"] for r in results):
        print(f"RESULT: SUPPORTED mean={mean_metric_value} std={std_metric_value} support_fraction={support_fraction}")
    elif any(not r["conjecture_holds"] for r in results):
        first_failing_seed = next(s for s, r in zip(seeds, results) if not r["conjecture_holds"])
        print(f"RESULT: FALSIFIED counterexample=\"\" first_failing_seed={first_failing_seed}")
    else:
        print("RESULT: INCONCLUSIVE no seeds tested")

```

6. Per-seed results

(no seed-level data — test crashed before producing TRIAL: lines)

7. Test stdout (last 2KB)

```
ture_holds': False, 'counterexample': ''}
TRIAL: {'metric_name': 'Hodge-theoretic Index vs Frege Proof Depth', 'metric_value': 1.0000000000000000}
TRIAL: {'metric_name': 'Hodge-theoretic Index vs Frege Proof Depth', 'metric_value': 1.0000000000000000}
TRIAL: {'metric_name': 'Hodge-theoretic Index vs Frege Proof Depth', 'metric_value': 1.0000000000000000}
TRIAL: {'metric_name': 'Hodge-theoretic Index vs Frege Proof Depth', 'metric_value': 1.0000000000000000}
TRIAL: {'metric_name': 'Hodge-theoretic Index vs Frege Proof Depth', 'metric_value': 1.0000000000000000}
TRIAL: {'metric_name': 'Hodge-theoretic Index vs Frege Proof Depth', 'metric_value': 1.0000000000000000}
TRIAL: {'metric_name': 'Hodge-theoretic Index vs Frege Proof Depth', 'metric_value': 1.0000000000000000}
TRIAL: {'metric_name': 'Hodge-theoretic Index vs Frege Proof Depth', 'metric_value': 1.0000000000000000}
TRIAL: {'metric_name': 'Hodge-theoretic Index vs Frege Proof Depth', 'metric_value': 1.0000000000000000}
TRIAL: {'metric_name': 'Hodge-theoretic Index vs Frege Proof Depth', 'metric_value': 1.0000000000000000}
TRIAL: {'metric_name': 'Hodge-theoretic Index vs Frege Proof Depth', 'metric_value': 1.0000000000000000}
RESULT: FALSIFIED counterexample="" first_failing_seed=11
```

8. Critic adversarial review

Critic verdict: CHALLENGE

Critic reasoning:

The test uses the number of clauses as a proxy for the Hodge-theoretic index, which is not correct. The Hodge-theoretic index is a complex algebraic geometric property that cannot be accurately represented by the count of clauses in a CNF. Additionally, the Frege proof depth calculation is simplified and does not reflect the actual complexity of proving a clause.

9. Final verdict & safety rail

Verdict: INCONCLUSIVE

Reasoning:

Safety rail: critic_challenge_falsified | original: The test results indicate that the conjecture does not hold for at least one seed (first_failing_seed=11), as $h(\varphi) < \Omega(d(\varphi))$ or $h(\varphi)/d(\varphi)^2 > c$ for so | next: Further investigation is needed to determine the exact relationship between Hodge-theoretic index and Frege proof depth. Consider using more accurate methods for calculating these metrics and retesting the conjecture.

11. Audit log (LLM calls)

Total LLM calls: 10

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
1	propose	ollama_remote	glm4:latest	0	0	16058
2	preregistration	ollama_remote	glm4:latest	0	0	10516
3	novelty	ollama_remote	glm4:latest	0	0	12265
4	novelty	ollama_remote	glm4:latest	0	0	8902
5	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	38651
6	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	10526
7	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	9254
8	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	18554
9	critic	ollama_remote	glm4:latest	0	0	50849

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
10	judge	ollama_remote	glm4:latest	0	0	10015

Totals: 0 input tokens, 0 output tokens, ~\$0.0000 cost, 185591 ms total latency. Provider mix: {'ollama_remote': 10}

(full prompt+response transcripts available in research/audit/10ae0a5c5991.jsonl)

12. Reproducibility

To reproduce this cycle:

1. Apply the test harness in section 5.1 with seeds 11,23,37,53,71
2. The Python sandbox is pure-stdlib; any Python 3.11+ should suffice
3. The full LLM transcripts are in research/audit/10ae0a5c5991.jsonl for verification of provenance
4. The pre-registration hash in section 2 commits to the test design before execution; tampering would be detectable.

Sandbox file: research/pvsnp_sandbox/test_*.py (per-cycle)

Replay tarball: research/replay/10ae0a5c5991.tar.gz (if generated)