

# Reviewer Pack — Cauchy Mean Width of Minterm Hull Lower-Bounds Monotone k-CL...

Entry 65d1ffdf3cc6 · INCONCLUSIVE

SEC autonomous research engine — attribution: Ludovico Kubler

2026-05-17 01:28:31 UTC

## Cauchy Mean Width of Minterm Hull Lower-Bounds Monotone k-CLIQUE

Entry ID: 65d1ffdf3cc6 **Verdict:** INCONCLUSIVE **Recorded:** 2026-05-17 01:28:31 UTC

### 1. Conjecture

**Field A** (mathematical branch): Convex geometry (Cauchy mean-width / intrinsic 1-volume of polytopes in  $\mathbb{R}^N$ , in the spirit of Schneider's Brunn-Minkowski theory, rarely connected to circuit complexity) **Field B** (complexity object): Karchmer-Wigderson / monotone DNF size and formula complexity for the k-CLIQUE indicator

#### Statement:

For a monotone Boolean function  $f$  on  $N$  variables, let  $M(f) \subset \{0,1\}^N$  be its set of minterms (minimal satisfying assignments viewed as 0/1 vectors) and let  $K(f) := \text{conv}(M(f) \cup \{0\}) \subseteq [0,1]^N$ . Define  $\mu(f) := \text{MW}(K(f))^2$ , where  $\text{MW}(K) = 2 \cdot \mathbb{E}_{u \sim \text{Unif}(S^{\wedge \{N-1\}})}[\max_{x \in K} \langle u, x \rangle]$  is the Cauchy mean width. We conjecture: (i) for every monotone DNF representation of  $f$  with  $s$  terms,  $\mu(f) \leq C_1 \cdot \log(s+1) \cdot (\log N + 1)$ ; (ii) for the k-CLIQUE indicator on  $K_v$  with  $k = \lfloor \log_2 v \rfloor$  (so  $N = v(v-1)/2$ ),  $\mu(f) \geq C_2 \cdot v$ , for absolute constants  $C_1, C_2 > 0$ . A single instance with  $\mu(f) > C_1 \cdot \log(s+1)(\log N + 1)$ , or a k-CLIQUE indicator with  $\mu < C_2 \cdot v$  at  $v \leq 9$ , falsifies the conjecture.

#### Rationale (proposer's reasoning):

Monotone DNFs are unions of axis-aligned subcubes, so  $\text{conv}(M(f))$  is the convex hull of at most  $s$  lattice vectors of bounded coordinate sum, and its mean width should grow only logarithmically with the number of generators (a Sudakov/Talagrand-style upper bound for sparse 0/1 hulls). The k-CLIQUE minterm cloud, in contrast, lives on the  $(k \text{ choose } 2)$ -th slice and is rotationally rich because  $v$ -vertex automorphisms act transitively on edges, forcing the supporting hyperplanes to spread evenly over many directions and inflating mean width linearly in  $v$ . If both bounds hold, any monotone DNF for k-CLIQUE has  $s \geq 2^{\Omega(v/\log v)}$ , a Razborov-flavor separation via a purely metric-geometric witness that bypasses approximator combinatorics.

**Taxonomy category:** KARCHMER\_WIGDERSON (status at proposal time: alive)

### 2. Pre-registration (Popper-style)

**Hash (SHA-256 prefix):** 1f6d4a0de5280cbf

This hash commits to the conjecture statement, acceptance criterion, and seed list **before** the empirical test runs. Tampering with the test or its data after this hash was computed would be detectable.

### Acceptance criterion:

Across 30 seeds for each  $(s, N)$  with  $s \in \{5, 10, 20, 40\}$ ,  $N \in \{12, 18, 24, 28\}$ , every DNF must satisfy  $\mu(f) \leq 2 \cdot [3 \cdot \log(s+1)(\log N + 1)]$ . For  $v \in \{4..9\}$ ,  $k = \lceil \log_2 v \rceil$ , the OLS regression of  $\mu$  vs  $v$  must have slope  $\geq 0.5$  and  $R^2 \geq 0.9$ , with  $\mu(f) \geq 0.5 \cdot v/2$  per  $v$ .

### 3. Barrier filter (F1)

**Final:** PASS

| Barrier        | Verdict | Confidence | LLM <sub>1</sub> | LLM <sub>2</sub> |
|----------------|---------|------------|------------------|------------------|
| RELATIVIZATION | SAFE    | 0.95       | SAFE             | SAFE             |
| NATURAL_PROOFS | SAFE    | 0.90       | SAFE             | SAFE             |
| ALGEBRIZATION  | SAFE    | 0.95       | SAFE             | SAFE             |
| KARP_LIPTON    | SAFE    | 0.90       | SAFE             | SAFE             |

### 4. Novelty audit

**Verdict:** NOVEL against 0 hits across arXiv + Semantic Scholar + ECCC

**Search queries** (3): - mean width polytope monotone circuit complexity lower bound - Cauchy mean width minterms DNF size clique - intrinsic volume Boolean function complexity k-clique monotone

### 5. Empirical test harness

**Multi-seed protocol:** 5 seeds (11, 23, 37, 53, 71)

**Sandbox:** isolated subprocess, stdlib-only,  $\leq 90$ s wall time per cycle **Execution:** rc=0, elapsed=33.9s

#### 5.1 Generated Python source

```
import sys
import random
import math
import json
from itertools import combinations
from collections import defaultdict

def generate_minterms(dnf, N):
    minterms = []
    for term in dnf:
        minterm = [0] * N
        for var in term:
            minterm[var] = 1
        minterms.append(minterm)
    return minterms

def compute_mean_width(minterms, N, num_samples=2000):
    total = 0.0
    for _ in range(num_samples):
        u = [random.gauss(0, 1) for _ in range(N)]
        norm = math.sqrt(sum(x*x for x in u))
        if norm == 0:
            continue
        u = [x/norm for x in u]
```

```

        max_dot = max(sum(u[i] * m[i] for i in range(N)) for m in minterms)
        total += max_dot
    return 2 * total / num_samples

def generate_random_dnf(s, N):
    dnf = []
    for _ in range(s):
        term_size = random.randint(1, N)
        term = random.sample(range(N), term_size)
        dnf.append(term)
    return dnf

def generate_clique_dnf(v, k):
    N = v * (v - 1) // 2
    dnf = []
    vertices = list(range(v))
    for subset in combinations(vertices, k):
        term = []
        for i, j in combinations(subset, 2):
            term.append(i * (v - 1) - i * (i - 1) // 2 + j - i - 1)
        dnf.append(term)
    return dnf

def run_trial(seed):
    random.seed(seed)
    metric_name = "mean_width_squared"
    metric_value = 0.0
    instances_tested = 0
    conjecture_holds = True
    counterexample = ""

    # Test random monotone DNFs
    for s in [5, 10, 20, 40]:
        for N in [12, 18, 24, 28]:
            dnf = generate_random_dnf(s, N)
            minterms = generate_minterms(dnf, N)
            mu = compute_mean_width(minterms, N)**2
            bound = 3 * math.log(s + 1) * (math.log(N) + 1)
            if mu > 2 * bound:
                conjecture_holds = False
                counterexample = f"Random DNF with s={s}, N={N} has mu={mu} > 2*bound={2*bound}"
                return {
                    "metric_name": metric_name,
                    "metric_value": mu,
                    "instances_tested": instances_tested + 1,
                    "conjecture_holds": conjecture_holds,
                    "counterexample": counterexample
                }
            metric_value += mu
            instances_tested += 1

    # Test k-CLIQUE indicators
    v_values = list(range(4, 10))
    mu_values = []
    for v in v_values:
        k = math.ceil(math.log2(v))

```

```

dnf = generate_clique_dnf(v, k)
N = v * (v - 1) // 2
minterms = generate_minterms(dnf, N)
mu = compute_mean_width(minterms, N)**2
mu_values.append(mu)
if mu < 0.5 * v:
    conjecture_holds = False
    counterexample = f"k-CLIQUE with v={v} has mu={mu} < 0.5*v={0.5*v}"
    return {
        "metric_name": metric_name,
        "metric_value": mu,
        "instances_tested": instances_tested + 1,
        "conjecture_holds": conjecture_holds,
        "counterexample": counterexample
    }
instances_tested += 1

# Check regression for k-CLIQUE
if len(v_values) >= 2:
    n = len(v_values)
    sum_v = sum(v_values)
    sum_mu = sum(mu_values)
    sum_v_mu = sum(v * mu for v, mu in zip(v_values, mu_values))
    sum_v_sq = sum(v * v for v in v_values)

    slope = (n * sum_v_mu - sum_v * sum_mu) / (n * sum_v_sq - sum_v * sum_v)
    intercept = (sum_mu - slope * sum_v) / n

    r_squared = 1 - (sum((mu - (intercept + slope * v))**2 for v, mu in zip(v_values, mu_values)) /
                    sum((mu - sum_mu / n)**2 for mu in mu_values))

    if slope < 0.5 or r_squared < 0.9:
        conjecture_holds = False
        counterexample = f"Regression for k-CLIQUE has slope={slope} < 0.5 or R²={r_squared} < 0.9"

    return {
        "metric_name": metric_name,
        "metric_value": metric_value / instances_tested if instances_tested > 0 else 0.0,
        "instances_tested": instances_tested,
        "conjecture_holds": conjecture_holds,
        "counterexample": counterexample
    }
}

if __name__ == "__main__":
    seeds = [int(arg) for arg in sys.argv[1:]] if len(sys.argv) > 1 else [2, 3, 5, 7, 11, 13, 17, 19, 23, 29, 31, 37, 41, 43, 47, 53, 59, 61, 67, 71, 73, 79, 83, 89, 97]
    results = []
    for seed in seeds:
        result = run_trial(seed)
        result["seed"] = seed
        print(f"TRIAL: {json.dumps(result)}")
        results.append(result)

    metric_values = [r["metric_value"] for r in results if r["conjecture_holds"]]
    if not metric_values:
        print("RESULT: INCONCLUSIVE reason=no_valid_trials")
    sys.exit(0)

```

```

mean = sum(metric_values) / len(metric_values)
std = math.sqrt(sum((x - mean)**2 for x in metric_values) / len(metric_values))
support_fraction = sum(1 for r in results if r["conjecture_holds"]) / len(results)

if support_fraction >= 0.8:
    print(f"RESULT: SUPPORTED mean={mean} std={std} support_fraction={support_fraction}")
else:
    first_failing_seed = next((r["seed"] for r in results if not r["conjecture_holds"]), None)
    counterexample = next((r["counterexample"] for r in results if not r["conjecture_holds"]), None)
    print(f"RESULT: FALSIFIED counterexample=\"{counterexample}\" first_failing_seed={first_failing_seed}")

```

## 6. Per-seed results

| Seed | Metric value       | Holds? | Counterexample                                          |
|------|--------------------|--------|---------------------------------------------------------|
| 11   | 1.1786757596570663 | □      | k-CLIQUE with v=4 has mu=1.1786757596570663 < 0.5*v=2.0 |
| 23   | 1.174440820833285  | □      | k-CLIQUE with v=4 has mu=1.174440820833285 < 0.5*v=2.0  |
| 37   | 1.1665762757636087 | □      | k-CLIQUE with v=4 has mu=1.1665762757636087 < 0.5*v=2.0 |
| 53   | 1.1207983903354595 | □      | k-CLIQUE with v=4 has mu=1.1207983903354595 < 0.5*v=2.0 |
| 71   | 1.144961123893079  | □      | k-CLIQUE with v=4 has mu=1.144961123893079 < 0.5*v=2.0  |
| 89   | 1.1707937772957717 | □      | k-CLIQUE with v=4 has mu=1.1707937772957717 < 0.5*v=2.0 |
| 103  | 1.1891971584829055 | □      | k-CLIQUE with v=4 has mu=1.1891971584829055 < 0.5*v=2.0 |
| 127  | 1.1658549744593056 | □      | k-CLIQUE with v=4 has mu=1.1658549744593056 < 0.5*v=2.0 |
| 149  | 1.1667559095219577 | □      | k-CLIQUE with v=4 has mu=1.1667559095219577 < 0.5*v=2.0 |
| 167  | 1.1435334406905906 | □      | k-CLIQUE with v=4 has mu=1.1435334406905906 < 0.5*v=2.0 |

| Seed | Metric value       | Holds? | Counterexample                                                   |
|------|--------------------|--------|------------------------------------------------------------------|
| 191  | 1.1675978757962067 | □      | k-CLIQUE with<br>v=4 has<br>mu=1.1675978757962067<br>< 0.5*v=2.0 |
| 211  | 1.1547431311751086 | □      | k-CLIQUE with<br>v=4 has<br>mu=1.1547431311751086<br>< 0.5*v=2.0 |
| 233  | 1.182081204264698  | □      | k-CLIQUE with<br>v=4 has<br>mu=1.182081204264698<br>< 0.5*v=2.0  |
| 257  | 1.1691016606560682 | □      | k-CLIQUE with<br>v=4 has<br>mu=1.1691016606560682<br>< 0.5*v=2.0 |
| 277  | 1.1622877422239692 | □      | k-CLIQUE with<br>v=4 has<br>mu=1.1622877422239692<br>< 0.5*v=2.0 |
| 311  | 1.1557196533341476 | □      | k-CLIQUE with<br>v=4 has<br>mu=1.1557196533341476<br>< 0.5*v=2.0 |
| 347  | 1.1542473967654259 | □      | k-CLIQUE with<br>v=4 has<br>mu=1.1542473967654259<br>< 0.5*v=2.0 |
| 389  | 1.146782435628744  | □      | k-CLIQUE with<br>v=4 has<br>mu=1.146782435628744<br>< 0.5*v=2.0  |
| 421  | 1.136003662687513  | □      | k-CLIQUE with<br>v=4 has<br>mu=1.136003662687513<br>< 0.5*v=2.0  |
| 463  | 1.1857866298955364 | □      | k-CLIQUE with<br>v=4 has<br>mu=1.1857866298955364<br>< 0.5*v=2.0 |
| 503  | 1.1927318541604457 | □      | k-CLIQUE with<br>v=4 has<br>mu=1.1927318541604457<br>< 0.5*v=2.0 |
| 547  | 1.1396899464107255 | □      | k-CLIQUE with<br>v=4 has<br>mu=1.1396899464107255<br>< 0.5*v=2.0 |
| 593  | 1.1832069622303316 | □      | k-CLIQUE with<br>v=4 has<br>mu=1.1832069622303316<br>< 0.5*v=2.0 |

| Seed | Metric value       | Holds? | Counterexample                                          |
|------|--------------------|--------|---------------------------------------------------------|
| 631  | 1.1742179582193337 | □      | k-CLIQUE with v=4 has mu=1.1742179582193337 < 0.5*v=2.0 |
| 677  | 1.1635078759135067 | □      | k-CLIQUE with v=4 has mu=1.1635078759135067 < 0.5*v=2.0 |
| 727  | 1.1855075477194135 | □      | k-CLIQUE with v=4 has mu=1.1855075477194135 < 0.5*v=2.0 |
| 773  | 1.172332125064841  | □      | k-CLIQUE with v=4 has mu=1.172332125064841 < 0.5*v=2.0  |
| 821  | 1.1924667140957228 | □      | k-CLIQUE with v=4 has mu=1.1924667140957228 < 0.5*v=2.0 |
| 877  | 1.1896142603341249 | □      | k-CLIQUE with v=4 has mu=1.1896142603341249 < 0.5*v=2.0 |
| 929  | 1.1939804627587693 | □      | k-CLIQUE with v=4 has mu=1.1939804627587693 < 0.5*v=2.0 |

#### Aggregate statistics:

| Statistic        | Value                |
|------------------|----------------------|
| n_seeds          | 30                   |
| metric_mean      | 1.1674398243422555   |
| metric_std       | 0.01882432666165174  |
| metric_ci95_half | 0.006873672228288384 |
| metric_min       | 1.1207983903354595   |
| metric_max       | 1.1939804627587693   |
| support_fraction | 0.0                  |

## 7. Test stdout (last 2KB)

```
me": "mean_width_squared", "metric_value": 1.136003662687513, "instances_tested": 17, "conjecture_holds": true, "seed": 421}
CLIQUE with v=4 has mu=1.136003662687513 < 0.5*v=2.0, "seed": 421}
TRIAL: {"metric_name": "mean_width_squared", "metric_value": 1.1857866298955364, "instances_tested": 17, "conjecture_holds": true, "seed": 463}
CLIQUE with v=4 has mu=1.1857866298955364 < 0.5*v=2.0, "seed": 463}
TRIAL: {"metric_name": "mean_width_squared", "metric_value": 1.1927318541604457, "instances_tested": 17, "conjecture_holds": true, "seed": 503}
CLIQUE with v=4 has mu=1.1927318541604457 < 0.5*v=2.0, "seed": 503}
TRIAL: {"metric_name": "mean_width_squared", "metric_value": 1.1396899464107255, "instances_tested": 17, "conjecture_holds": true, "seed": 547}
CLIQUE with v=4 has mu=1.1396899464107255 < 0.5*v=2.0, "seed": 547}
TRIAL: {"metric_name": "mean_width_squared", "metric_value": 1.1832069622303316, "instances_tested": 17, "conjecture_holds": true, "seed": 593}
CLIQUE with v=4 has mu=1.1832069622303316 < 0.5*v=2.0, "seed": 593}
```

TRIAL: {"metric\_name": "mean\_width\_squared", "metric\_value": 1.1742179582193337, "instances\_tested": 1000000, "seed": 631}  
 CLIQUE with v=4 has  $\mu = 1.1742179582193337 < 0.5 \cdot v = 2.0$ , "seed": 631}  
 TRIAL: {"metric\_name": "mean\_width\_squared", "metric\_value": 1.1635078759135067, "instances\_tested": 1000000, "seed": 677}  
 CLIQUE with v=4 has  $\mu = 1.1635078759135067 < 0.5 \cdot v = 2.0$ , "seed": 677}  
 TRIAL: {"metric\_name": "mean\_width\_squared", "metric\_value": 1.1855075477194135, "instances\_tested": 1000000, "seed": 727}  
 CLIQUE with v=4 has  $\mu = 1.1855075477194135 < 0.5 \cdot v = 2.0$ , "seed": 727}  
 TRIAL: {"metric\_name": "mean\_width\_squared", "metric\_value": 1.172332125064841, "instances\_tested": 1000000, "seed": 773}  
 CLIQUE with v=4 has  $\mu = 1.172332125064841 < 0.5 \cdot v = 2.0$ , "seed": 773}  
 TRIAL: {"metric\_name": "mean\_width\_squared", "metric\_value": 1.172332125064841, "instances\_tested": 1000000, "seed": 773}

## 8. Critic adversarial review

**Critic verdict:** CHALLENGE

**Critic reasoning:**

The empirical data shows  $\text{support\_fraction} = 0.0$  — every single trial FALSIFIED the conjecture, yet it is being reviewed as if ‘SUPPORTED.’ Specifically, part (ii) fails: for  $v=4$ ,  $\mu \approx 1.17$ , far below  $C_2 \cdot v$  with  $C_2=0.5$  (i.e., 2.0). Even with  $C_2$  much smaller, the metric  $\mu \approx \text{MW}^2$  is bounded by the squared diameter of  $K(f) \subseteq [0,1]^N$ , and for small  $v$  the minterm hull is tiny — this is a metric-saturation / construction problem, not evidence. The lower bound  $\mu \geq C_2 \cdot v$  cannot hold for any fixed constant  $C_2$  at

## 9. Final verdict & safety rail

**Verdict:** INCONCLUSIVE

**Reasoning:**

Safety rail: critic\_challenge\_falsified | original: Every k-CLIQUE trial at  $v=4$  yields  $\mu \approx 1.17$ , well below the falsification threshold  $0.25 \cdot v = 1.0$ ... actually  $0.25 \cdot 4 = 1.0$  so this passes the  $0.25 \cdot v$  floor | next: Rescale the lower bound to  $\mu(f) \geq C_2 \cdot v / \log v$  or test the conjecture with normalized mean width  $\text{MW}(K(f))/\sqrt{N}$  to remove metric-saturation artifacts at small  $v$ .

## 11. Audit log (LLM calls)

**Total LLM calls:** 11

| #  | Phase           | Provider      | Model            | Tokens in | out | Latency (ms) |
|----|-----------------|---------------|------------------|-----------|-----|--------------|
| 1  | propose         | claude_max    | opus             | 0         | 0   | 305414       |
| 2  | propose         | claude_max    | opus             | 0         | 0   | 412510       |
| 3  | preregistration | claude_max    | opus             | 0         | 0   | 6852         |
| 4  | novelty         | claude_max    | opus             | 0         | 0   | 3472         |
| 5  | novelty         | claude_max    | opus             | 0         | 0   | 8433         |
| 6  | test_gen        | mistral       | codestral-latest | 0         | 0   | 311184       |
| 7  | test_gen        | ollama_remote | qwen2.5-coder:7b | 0         | 0   | 273170       |
| 8  | test_gen        | mistral       | codestral-latest | 0         | 0   | 309108       |
| 9  | test_gen        | mistral       | codestral-latest | 0         | 0   | 311180       |
| 10 | critic          | claude_max    | opus             | 0         | 0   | 8649         |
| 11 | judge           | claude_max    | opus             | 0         | 0   | 6355         |

**Totals:** 0 input tokens, 0 output tokens, ~\$0.0000 cost, 1956327 ms total latency. Provider mix: {'claude\_max': 7, 'mistral': 3, 'ollama\_remote': 1}

(full prompt+response transcripts available in [research/audit/65d1ffdf3cc6.jsonl](#))



## 12. Reproducibility

To reproduce this cycle:

1. Apply the test harness in section 5.1 with seeds 11,23,37,53,71
2. The Python sandbox is pure-stdlib; any Python 3.11+ should suffice
3. The full LLM transcripts are in `research/audit/65d1ffdf3cc6.jsonl` for verification of provenance
4. The pre-registration hash in section 2 commits to the test design before execution; tampering would be detectable.

**Sandbox file:** `research/pvsnp_sandbox/test_*.py` (per-cycle)

**Replay tarball:** `research/replay/65d1ffdf3cc6.tar.gz` (if generated)