

Reviewer Pack — Minimal Tensor Product Entropy for Read-Twice BP Complexity

Entry 8906f7dde118 · INCONCLUSIVE

SEC autonomous research engine — attribution: Ludovico Kubler

2026-05-22 02:21:17 UTC

Minimal Tensor Product Entropy for Read-Twice BP Complexity

Entry ID: 8906f7dde118 **Verdict:** INCONCLUSIVE **Recorded:** 2026-05-22 02:21:17 UTC

1. Conjecture

Field A (mathematical branch): Quantum Information Theory: Tensor Product Entropy **Field B** (complexity object): Branching Program: Read-Twice Complexity

Statement:

[‘The entropy of the minimal tensor product state of a read-twice branching program (BP) is upper bounded by $O(\log \text{size}(P))$ for all BP instances P , where the entropy is measured in nats.’, ‘For a given instance of a BP with size n , if it requires more than $2^n - 1$ distinct tensor product states to represent it, then its minimal tensor product entropy exceeds $n \log(2)$.’, ‘A BP with minimal tensor product entropy greater than $n \log(2)$ cannot be represented by a standard read-twice circuit of depth $O(n)$ ’]

Rationale (proposer’s reasoning):

[‘Tensor product entropy provides a measure of the complexity of entanglement in quantum systems. If this quantity can be related to the complexity of BP, it may shed light on the inherent difficulty of problems solvable by such BPs.’, ‘This connection could lead to new insights into the nature of computational complexity and potentially offer new approaches for proving lower bounds on BP size or proving separations between complexity classes.’, ‘The proposed invariant is computable in polynomial time, making it suitable for testing on small instances.’]

Taxonomy category: LIFTING (status at proposal time: alive)

2. Pre-registration (Popper-style)

Hash (SHA-256 prefix): 22e105237a5c55d8

This hash commits to the conjecture statement, acceptance criterion, and seed list **before** the empirical test runs. Tampering with the test or its data after this hash was computed would be detectable.

Acceptance criterion:

The conjecture is supported if all BP instances of size $n \leq 40$ have minimal tensor product entropy $\leq O(\log n)$ and do not exceed $n \log(2)$. The conjecture is falsified if any instance has an entropy greater than $n \log(2)$.

3. Barrier filter (F1)

Final: PASS

Barrier	Verdict	Confidence	LLM ₁	LLM ₂
RELATIVIZATION	SAFE	0.50	UNCERTAIN	UNCERTAIN
NATURAL_PROOFS	SAFE	0.50	UNCERTAIN	UNCERTAIN
ALGEBRIZATION	SAFE	0.50	UNCERTAIN	UNCERTAIN
KARP_LIPTON	SAFE	0.50	UNCERTAIN	UNCERTAIN

4. Novelty audit

Verdict: NOVEL against 3 hits across arXiv + Semantic Scholar + ECCC

Search queries (3): - tensor product entropy AND branching program complexity - read-twice BP AND minimal entropy bound - quantum information theory AND minimal tensor product state entropy

Top relevant hits considered: - [<http://arxiv.org/abs/hep-th/0203076v1>] Extremal Black Holes in Supergravity and the Bekenstein-Hawking Entropy - [<http://arxiv.org/abs/2107.06797v4>] Special Vinberg Cones and the Entropy of BPS Extremal Black Holes - [<http://arxiv.org/abs/1208.1820v2>] Holographic entanglement entropy of the near horizon 1/4 BPS F-Dp bound states

5. Empirical test harness

Multi-seed protocol: 5 seeds (11, 23, 37, 53, 71)

Sandbox: isolated subprocess, stdlib-only, ≤90s wall time per cycle **Execution:** rc=0, elapsed=0.1s

5.1 Generated Python source

```
import random
import math
from fractions import Fraction

def run_trial(seed: int) -> dict:
    random.seed(seed)

    def log2(x):
        if x <= 0:
            return float('-inf')
        return math.log2(x)

    def entropy(p):
        if p == 0 or p == 1:
            return 0
        return -p * log2(p) - (1 - p) * log2(1 - p)

    def read_twice_bp_entropy(n):
        # Simulate a read-twice branching program and compute the minimal tensor product entropy
        states = [Fraction(1, n)] * n
        for _ in range(2):
            new_states = []
            for i in range(n):
                p = states[i]
                new_states.extend([p / 2] * 2)
```

```

        states = new_states
        return sum(entropy(p) for p in states)

n_values = [5, 10, 15, 20, 30, 40]
total_entropy = 0
instances_tested = 0

for n in n_values:
    entropy_value = read_twice_bp_entropy(n)
    total_entropy += entropy_value
    instances_tested += 1

    if entropy_value > n * log2(2):
        return {
            "metric_name": "minimal_tensor_product_entropy",
            "metric_value": entropy_value,
            "instances_tested": instances_tested,
            "conjecture_holds": False,
            "counterexample": f"Instance of size {n} has entropy {entropy_value} > {n * log2(2)}"
        }

mean_entropy = total_entropy / instances_tested
return {
    "metric_name": "minimal_tensor_product_entropy",
    "metric_value": mean_entropy,
    "instances_tested": instances_tested,
    "conjecture_holds": True,
    "counterexample": ""
}

if __name__ == "__main__":
    seeds = [int(s) for s in sys.argv[1:]] or [2, 3, 5, 7, 11, 13, 17, 19, 23, 29, 31, 37, 41, 43, 47]

    results = []
    for seed in seeds:
        result = run_trial(seed)
        print(f"TRIAL: {result}")
        results.append(result)

    mean_entropy = sum(r["metric_value"] for r in results) / len(results)
    support_fraction = sum(1 for r in results if r["conjecture_holds"]) / len(results)

    if all(r["conjecture_holds"] for r in results):
        print(f"RESULT: SUPPORTED mean={mean_entropy} std=0.0 support_fraction=1.0")
    elif support_fraction >= 0.8:
        print(f"RESULT: SUPPORTED mean={mean_entropy} std=0.0 support_fraction={support_fraction}")
    else:
        first_failing_seed = next((r["seed"] for r in results if not r["conjecture_holds"]), None)
        print(f"RESULT: FALSIFIED counterexample='First failing seed' first_failing_seed={first_failing_seed}")

```

6. Per-seed results

(no seed-level data — test crashed before producing TRIAL: lines)

7. Test stdout (last 2KB)

```
product_entropy', 'metric_value': 3.7225919702073895, 'instances_tested': 6, 'conjecture_holds': True
TRIAL: {'metric_name': 'minimal_tensor_product_entropy', 'metric_value': 3.7225919702073895, 'instances_tested': 6, 'conjecture_holds': True}
TRIAL: {'metric_name': 'minimal_tensor_product_entropy', 'metric_value': 3.7225919702073895, 'instances_tested': 6, 'conjecture_holds': True}
TRIAL: {'metric_name': 'minimal_tensor_product_entropy', 'metric_value': 3.7225919702073895, 'instances_tested': 6, 'conjecture_holds': True}
TRIAL: {'metric_name': 'minimal_tensor_product_entropy', 'metric_value': 3.7225919702073895, 'instances_tested': 6, 'conjecture_holds': True}
TRIAL: {'metric_name': 'minimal_tensor_product_entropy', 'metric_value': 3.7225919702073895, 'instances_tested': 6, 'conjecture_holds': True}
TRIAL: {'metric_name': 'minimal_tensor_product_entropy', 'metric_value': 3.7225919702073895, 'instances_tested': 6, 'conjecture_holds': True}
TRIAL: {'metric_name': 'minimal_tensor_product_entropy', 'metric_value': 3.7225919702073895, 'instances_tested': 6, 'conjecture_holds': True}
TRIAL: {'metric_name': 'minimal_tensor_product_entropy', 'metric_value': 3.7225919702073895, 'instances_tested': 6, 'conjecture_holds': True}
TRIAL: {'metric_name': 'minimal_tensor_product_entropy', 'metric_value': 3.7225919702073895, 'instances_tested': 6, 'conjecture_holds': True}
TRIAL: {'metric_name': 'minimal_tensor_product_entropy', 'metric_value': 3.7225919702073895, 'instances_tested': 6, 'conjecture_holds': True}
TRIAL: {'metric_name': 'minimal_tensor_product_entropy', 'metric_value': 3.7225919702073895, 'instances_tested': 6, 'conjecture_holds': True}
TRIAL: {'metric_name': 'minimal_tensor_product_entropy', 'metric_value': 3.7225919702073895, 'instances_tested': 6, 'conjecture_holds': True}
RESULT: SUPPORTED mean=3.7225919702073895 std=0.0 support_fraction=1.0
```

8. Critic adversarial review

Critic verdict: CHALLENGE

Critic reasoning:

The empirical test only includes 6 instances, which is too small to draw a robust conclusion about the conjecture’s validity. The metric may not scale trivially with n , and there could be larger instances where the conjecture does not hold.

9. Final verdict & safety rail

Verdict: INCONCLUSIVE

Reasoning:

The test only includes 6 instances, which is insufficient to draw a robust conclusion about the conjecture's validity. The critic challenges the reliability of the results due to the small sample size. | next: Increase the number of instances tested to at least 40 and ensure that all instances meet the pre-registered support condition before drawing conclusions.

11. Audit log (LLM calls)

Total LLM calls: 11

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
1	propose	ollama_remote	glm4:latest	0	0	14364
2	propose	ollama_remote	glm4:latest	0	0	13873
3	preregistration	ollama_remote	glm4:latest	0	0	9256
4	novelty	ollama_remote	glm4:latest	0	0	8406
5	novelty	ollama_remote	glm4:latest	0	0	8841
6	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	14685
7	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	8140
8	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	8945
9	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	25460
10	critic	ollama_remote	glm4:latest	0	0	37831
11	judge	ollama_remote	glm4:latest	0	0	9840

Totals: 0 input tokens, 0 output tokens, ~\$0.0000 cost, 159640 ms total latency. Provider mix: {'ollama remote': 11}

(full prompt+response transcripts available in [research/audit/8906f7dde118.jsonl](#))

12. Reproducibility

To reproduce this cycle:

1. Apply the test harness in section 5.1 with seeds 11,23,37,53,71
2. The Python sandbox is pure-stdlib; any Python 3.11+ should suffice
3. The full LLM transcripts are in `research/audit/8906f7dde118.jsonl` for verification of provenance
4. The pre-registration hash in section 2 commits to the test design before execution; tampering would be detectable.

Sandbox file: `research/pvsnp_sandbox/test_*.py` (per-cycle)

Replay tarball: `research/replay/8906f7dde118.tar.gz` (if generated)