

Reviewer Pack — Minimal Order of Automorphism Groups and Communication Compl...

Entry 47419d3f094e · INCONCLUSIVE

SEC autonomous research engine — attribution: Ludovico Kubler

2026-05-31 12:40:42 UTC

Minimal Order of Automorphism Groups and Communication Complexity via Noncrossing Partition Complexity

Entry ID: 47419d3f094e **Verdict:** INCONCLUSIVE **Recorded:** 2026-05-31 12:40:42 UTC

1. Conjecture

Field A (mathematical branch): Combinatorial Enumeration **Field B** (complexity object): Communication Complexity

Statement:

The communication complexity of a function f over an n -dimensional domain is linearly correlated with the minimal order of automorphism groups acting on its noncrossing partition complex, such that $|\Omega(f)| = \Theta(O(f))$.

Rationale (proposer's reasoning):

This conjecture bridges combinatorial enumeration theory with communication complexity by suggesting that the symmetry of a function can be captured through its noncrossing partitions, potentially revealing new insights into the hardness of communication tasks.

Taxonomy category: Combinatorial_Enumeration (status at proposal time:)

2. Pre-registration (Popper-style)

Hash (SHA-256 prefix): a77d045f7df2904e

This hash commits to the conjecture statement, acceptance criterion, and seed list **before** the empirical test runs. Tampering with the test or its data after this hash was computed would be detectable.

Acceptance criterion:

The conjecture is supported if, across all 30 random seeds, the ratio of the minimal order of automorphism groups $|\Omega(f)|$ to the communication complexity $O(f)$ for each function f is within a margin of error of 1 with a standard deviation less than 0.5.

3. Barrier filter (F1)

Final: PASS

Barrier	Verdict	Confidence	LLM ₁	LLM ₂
RELATIVIZATION	SAFE	1.00	UNCERTAIN	SAFE

Barrier	Verdict	Confidence	LLM ₁	LLM ₂
NATURAL_PROOFS	SAFE	0.50	UNCERTAIN	UNCERTAIN
ALGEBRIZATION	SAFE	0.90	SAFE	UNCERTAIN
KARP_LIPTON	SAFE	0.50	UNCERTAIN	UNCERTAIN

4. Novelty audit

Verdict: NOVEL against 0 hits across arXiv + Semantic Scholar + ECCC

5. Empirical test harness

Multi-seed protocol: 5 seeds (11, 23, 37, 53, 71)

Sandbox: isolated subprocess, stdlib-only, ≤90s wall time per cycle **Execution:** rc=124, elapsed=240.6s

5.1 Generated Python source

```
import random
import math

def run_trial(seed: int) -> dict:
    random.seed(seed)

    def generate_random_boolean_function(n):
        return [random.choice([0, 1]) for _ in range(2**n)]

    def noncrossing_partition_complex(f, n):
        # Placeholder function to compute the noncrossing partition complex
        # This is a dummy implementation and should be replaced with actual logic
        return f

    def automorphism_group_order(noncrossing_partition):
        # Placeholder function to compute the order of the automorphism group
        # This is a dummy implementation and should be replaced with actual logic
        return len(noncrossing_partition)

    def communication_complexity(f, n):
        # Placeholder function to compute the communication complexity
        # This is a dummy implementation and should be replaced with actual logic
        return n

    metric_name = "communication_complexity"
    instances_tested = 0
    total_metric_value = 0.0
    n_max = 1
    conjecture_holds = True
    counterexample = ""

    for _ in range(30):
        n = random.choice([5, 10, 15, 20, 30, 40])
        if n > n_max:
            n_max = n

        f = generate_random_boolean_function(n)
        noncrossing_partition = noncrossing_partition_complex(f, n)
```

```

order_of_automorphism_group = automorphism_group_order(noncrossing_partition)
comm_complexity = communication_complexity(f, n)

if comm_complexity == 0:
    continue

instances_tested += 1
total_metric_value += order_of_automorphism_group / comm_complexity

if instances_tested > 0:
    mean_C = total_metric_value / instances_tested
    std_dev = math.sqrt(sum((order_of_automorphism_group / comm_complexity - mean_C) ** 2 for
                             _ in range(instances_tested)))

    if mean_C < 1 or std_dev >= 0.5:
        conjecture_holds = False
        counterexample = "mean_C out of bounds"

return {
    "metric_name": metric_name,
    "metric_value": mean_C,
    "instances_tested": instances_tested,
    "n_max": n_max,
    "conjecture_holds": conjecture_holds,
    "counterexample": counterexample
}

if __name__ == "__main__":
    seeds = [int(x) for x in sys.argv[1:]] or [2, 3, 5, 7, 11, 13, 17, 19, 23, 29, 31, 37, 41, 43, 47]

    results = []
    for seed in seeds:
        result = run_trial(seed)
        print(f"TRIAL: {{\"seed\": {seed}, **result}}")
        results.append(result)

    if all(result["conjecture_holds"] for result in results):
        mean_C = sum(result["metric_value"] for result in results) / len(results)
        std_dev = math.sqrt(sum((result["metric_value"] - mean_C) ** 2 for result in results) / len(results))
        print(f"RESULT: SUPPORTED mean={mean_C} std={std_dev} support_fraction=1.0")
    elif any(not result["conjecture_holds"] for result in results):
        first_failing_seed = next(seed for seed, result in zip(seeds, results) if not result["conjecture_holds"])
        print(f"RESULT: FALSIFIED counterexample=\"mean_C out of bounds\" first_failing_seed={first_failing_seed}")
    else:
        print("RESULT: INCONCLUSIVE reason=unreachable")

```

6. Per-seed results

(no seed-level data — test crashed before producing TRIAL: lines)

7. Test stdout (last 2KB)

TIMEOUT after 240s

8. Critic adversarial review

Critic verdict: CHALLENGE

Critic reasoning:

skipped: test crashed before producing data

9. Final verdict & safety rail

Verdict: INCONCLUSIVE

Reasoning:

The test timed out before producing data, which means it did not complete within the expected time frame. | next: NONE

11. Audit log (LLM calls)

Total LLM calls: 9

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
1	propose	ollama_remote	glm4:latest	0	0	17754
2	preregistration	ollama_remote	glm4:latest	0	0	14991
3	novelty	ollama_remote	glm4:latest	0	0	15702
4	novelty	ollama_remote	glm4:latest	0	0	21926
5	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	11670
6	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	19100
7	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	11859
8	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	16632
9	judge	ollama_remote	glm4:latest	0	0	32056

Totals: 0 input tokens, 0 output tokens, ~\$0.0000 cost, 161691 ms total latency. Provider mix: {'ollama_remote': 9}

(full prompt+response transcripts available in `research/audit/47419d3f094e.jsonl`)

12. Reproducibility

To reproduce this cycle:

1. Apply the test harness in section 5.1 with seeds 11,23,37,53,71
2. The Python sandbox is pure-stdlib; any Python 3.11+ should suffice
3. The full LLM transcripts are in `research/audit/47419d3f094e.jsonl` for verification of provenance
4. The pre-registration hash in section 2 commits to the test design before execution; tampering would be detectable.

Sandbox file: `research/pvsnp_sandbox/test_*.py` (per-cycle)

Replay tarball: `research/replay/47419d3f094e.tar.gz` (if generated)