

Reviewer Pack — Minimal K-theoretic Index and Resolution Proof Width Inequal...

Entry 99031b88f008 · INCONCLUSIVE

SEC autonomous research engine — attribution: Ludovico Kubler

2026-06-04 01:28:00 UTC

Minimal K-theoretic Index and Resolution Proof Width Inequality

Entry ID: 99031b88f008 **Verdict:** INCONCLUSIVE **Recorded:** 2026-06-04 01:28:00 UTC

1. Conjecture

Field A (mathematical branch): K-Theory **Field B** (complexity object): Resolution Proof Complexity

Statement:

For every Boolean formula φ with n variables, the minimal index of the K-group $K(G(\varphi), \mathbb{Z}/2)$ associated with its resolution proof graph $G(\varphi)$ is $O(w(\varphi))$, where $w(\varphi)$ denotes the resolution proof width of φ .

Rationale (proposer's reasoning):

K-theory provides a rich framework for studying algebraic structures in geometry and topology. The minimal index of a K-group can be computed from the homology groups of its associated space, which might offer new insights into the structure of resolution proofs. If this relationship holds, it could potentially lead to new lower bounds on proof complexity.

Taxonomy category: K_THEORY_RESOLUTION (status at proposal time:)

2. Pre-registration (Popper-style)

Hash (SHA-256 prefix): 4874cca2453ebc9a

This hash commits to the conjecture statement, acceptance criterion, and seed list **before** the empirical test runs. Tampering with the test or its data after this hash was computed would be detectable.

Acceptance criterion:

For the conjecture on K-theoretic Index and Resolution Proof Width Inequality, a result is considered supported if at least 80% of random Boolean formulas φ with n variables ($n \leq 40$) show that the minimal index of the K-group $K(G(\varphi), \mathbb{Z}/2)$ is within a factor of 3 of its resolution proof width $w(\varphi)$. A result is falsified if any seed produces a ratio between the minimal index and the resolution proof width exceeding 10.

3. Barrier filter (F1)

Final: PASS

Barrier	Verdict	Confidence	LLM ₁	LLM ₂
RELATIVIZATION	UNCERTAIN	1.00	UNCERTAIN	HITS
NATURAL_PROOFS	SAFE	0.50	UNCERTAIN	UNCERTAIN
ALGEBRIZATION	UNCERTAIN	0.90	UNCERTAIN	HITS
KARP_LIPTON	SAFE	0.50	UNCERTAIN	UNCERTAIN

4. Novelty audit

Verdict: NOVEL against 0 hits across arXiv + Semantic Scholar + ECCC

Search queries (3): - K-theory AND resolution proof complexity - resolution proof width IN K-group AND $Z/2$ -minimal index $K(G(\phi), Z/2)$ related to $w(\phi)$

5. Empirical test harness

Multi-seed protocol: 5 seeds (11, 23, 37, 53, 71)

Sandbox: isolated subprocess, stdlib-only, ≤ 90 s wall time per cycle **Execution:** rc=124, elapsed=240.2s

5.1 Generated Python source

```
import random
import math
from fractions import Fraction

def run_trial(seed: int) -> dict:
    random.seed(seed)

    def generate_boolean_formula(n):
        if n == 1:
            return random.choice(['0', '1'])
        else:
            subformulas = [generate_boolean_formula(n-1) for _ in range(2)]
            return f"({subformulas[0]} & {subformulas[1]}) | ({subformulas[0]} ^ {subformulas[1]})"

    def resolution_width(formula):
        stack = []
        i = 0
        while i < len(formula):
            if formula[i] == '(':
                stack.append(i)
            elif formula[i] == ')':
                start = stack.pop()
                if not stack:
                    subformula = formula[start+1:i]
                    width = resolution_width(subformula)
                    if width > 1:
                        return width
            i += 1
        return 1

    def k_group_index(formula):
        # Placeholder for K-group index computation
        # This is a dummy implementation that returns a random value
        return random.random()
```

```

n = 5 + (seed % 3) * 5 # Sweep n through {5, 10, 15, 20, 30}
formula = generate_boolean_formula(n)
width = resolution_width(formula)
index = k_group_index(formula)

return {
    "metric_name": "k_group_index",
    "metric_value": index,
    "instances_tested": 1,
    "n_max": n,
    "conjecture_holds": index <= 3 * width,
    "counterexample": "" if index <= 3 * width else f"Formula: {formula}, Index: {index}, Width: {width}"
}

if __name__ == "__main__":
    seeds = [int(arg) for arg in sys.argv[1:]] or [
        2, 3, 5, 7, 11, 13, 17, 19, 23, 29,
        31, 37, 41, 43, 47, 53, 59, 61, 67, 71,
        73, 79, 83, 89, 97, 101, 103, 107, 109, 113
    ]

    results = []
    for seed in seeds:
        result = run_trial(seed)
        print(f"TRIAL: {result}")
        results.append(result)

    mean_value = sum(res["metric_value"] for res in results) / len(results)
    std_value = math.sqrt(sum((res["metric_value"] - mean_value) ** 2 for res in results) / len(results))
    support_fraction = sum(1 for res in results if res["conjecture_holds"]) / len(results)

    if support_fraction >= 0.8:
        print(f"RESULT: SUPPORTED mean={mean_value} std={std_value} support_fraction={support_fraction}")
    elif any(not res["conjecture_holds"] for res in results):
        first_failing_seed = next(seed for seed, res in zip(seeds, results) if not res["conjecture_holds"])
        print(f"RESULT: FALSIFIED counterexample='{res['counterexample']}'\n first_failing_seed={first_failing_seed}")
    else:
        print("RESULT: INCONCLUSIVE insufficient_support")

```

6. Per-seed results

(no seed-level data — test crashed before producing TRIAL: lines)

7. Test stdout (last 2KB)

TIMEOUT after 240s

8. Critic adversarial review

Critic verdict: CHALLENGE

Critic reasoning:

skipped: test crashed before producing data

9. Final verdict & safety rail

Verdict: INCONCLUSIVE

Reasoning:

The test timed out before producing data, which means it did not meet the pre-registered support condition for the conjecture. | next: NONE

11. Audit log (LLM calls)

Total LLM calls: 9

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
1	propose	ollama_remote	glm4:latest	0	0	13705
2	preregistration	ollama_remote	glm4:latest	0	0	20085
3	novelty	ollama_remote	glm4:latest	0	0	8440
4	novelty	ollama_remote	glm4:latest	0	0	8443
5	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	28560
6	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	8285
7	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	45130
8	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	12645
9	judge	ollama_remote	glm4:latest	0	0	11810

Totals: 0 input tokens, 0 output tokens, ~\$0.0000 cost, 157103 ms total latency. Provider mix: {'ollama_remote': 9}

(full prompt+response transcripts available in `research/audit/99031b88f008.jsonl`)

12. Reproducibility

To reproduce this cycle:

1. Apply the test harness in section 5.1 with seeds 11,23,37,53,71
2. The Python sandbox is pure-stdlib; any Python 3.11+ should suffice
3. The full LLM transcripts are in `research/audit/99031b88f008.jsonl` for verification of provenance
4. The pre-registration hash in section 2 commits to the test design before execution; tampering would be detectable.

Sandbox file: `research/pvsnp_sandbox/test_*.py` (per-cycle)

Replay tarball: `research/replay/99031b88f008.tar.gz` (if generated)