Reviewer Pack — Minimal Rank of Formal Power Series Bounds Resolution Proof ...

Entry 988964c281c1 · INCONCLUSIVE

SEC autonomous research engine — attribution: Ludovico Kubler

2026-05-29 14:22:25 UTC

Minimal Rank of Formal Power Series Bounds Resolution Proof Entropy

Entry ID: `988964c281c1` **Verdict**: `INCONCLUSIVE` **Recorded**: 2026-05-29 14:22:25 UTC

1. Conjecture

Field A (mathematical branch): Formal Power Series Theory **Field B** (complexity object): Resolution proof complexity

Statement:

> For a given Tseitin formula φ with n variables and m clauses, the minimal rank r of its associated formal power series F(φ) is bounded by O(m^(1/3)) in terms of resolution proof entropy Φ(φ), i.e., |χ_dir − B(φ)| ≤ k for some constant k.

Rationale (proposer's reasoning):

> Formal power series theory provides a framework to encode the structure of Boolean functions, and its minimal rank could potentially reveal hidden complexity properties. This conjecture suggests that the geometric complexity of the formal power series is related to the resolution proof entropy, which measures the difficulty of proving a formula in resolution.

Taxonomy category: `c003b_cumulative_entropy` (status at proposal time:)

2. Pre-registration (Popper-style)

Hash (SHA-256 prefix): `d9b7ced607bcc393`

This hash commits to the conjecture statement, acceptance criterion, and seed list **before** the empirical test runs. Tampering with the test or its data after this hash was computed would be detectable.

Acceptance criterion:

> The conjecture is supported if, for all generated Tseitin formulas with n variables and m clauses (n ≤ 40), the minimal rank r of F(φ) is bounded by O(m^(1/3)) in terms of resolution proof entropy Φ(φ), with no seed producing a metric |χ_dir − B(φ)| > k.

3. Barrier filter (F1)

Final: `PASS`

Barrier	Verdict	Confidence	LLM_1	LLM_2
RELATIVIZATION	SAFE	1.00	SAFE	SAFE
NATURAL_PROOFS	SAFE	0.50	UNCERTAIN	UNCERTAIN
ALGEBRIZATION	SAFE	0.50	UNCERTAIN	UNCERTAIN
KARP_LIPTON	SAFE	1.00	SAFE	SAFE

4. Novelty audit

Verdict: `NOVEL` against 0 hits across arXiv + Semantic Scholar + ECCC

Search queries (3): - `formal power series theory AND resolution proof complexity` - `minimal rank formal power series resolution proof entropy` - `Tseitin formula formal power series O(m^(1/3)) resolution proof entropy`

5. Empirical test harness

Multi-seed protocol: 5 seeds (11, 23, 37, 53, 71)
Sandbox: isolated subprocess, stdlib-only, ≤90s wall time per cycle **Execution**: rc=0, elapsed=0.4s

5.1 Generated Python source

```python
#!/usr/bin/env python3
"""C-003b cumulative entropy — FAITHFUL reference harness.

Computes the REAL cumulative proof entropy
    Phi(pi) = sum_t |activeClauses(sigma_t)|
over an actual resolution refutation (Davis-Putnam variable elimination)
of Tseitin formulas, exactly per the Lean definitions in Conjecture003.lean:
    proofStateEntropy(sigma) := sigma.activeClauses.card
    cumulativeEntropy(states) := sum (map proofStateEntropy states)
plus the width-aware totalLiteralWeight version (sum of clause sizes).

This is NOT the CDCL clause-count proxy used by the old c003b_counterexample.py:
it tracks the active clause database at EVERY elimination step of a genuine
resolution derivation, so it is the actual Phi the formalization names.

Pure stdlib (no networkx/pysat), so it runs in the pipeline sandbox.
Protocol: run_trial(seed)->dict with TRIAL/RESULT lines.

Author: Ludovico Kubler (harness by the SEC engine, hand-verified).
NOTE: no `from __future__ import annotations` — the sandbox injects a prelude
above the file, which would break that statement's must-be-first rule.
"""
import json
import math
import random
import sys

Clause = frozenset  # a clause = frozenset of signed ints (var or -var)


# ---------- Tseitin formula generation ----------

def random_regular_graph(n: int, degree: int, rng: random.Random):
```

```python
    """Simple random regular graph via configuration model with retries."""
    if (n * degree) % 2 != 0:
        n += 1
    for _ in range(200):
        stubs = []
        for v in range(n):
            stubs += [v] * degree
        rng.shuffle(stubs)
        edges = set()
        ok = True
        for i in range(0, len(stubs), 2):
            u, w = stubs[i], stubs[i + 1]
            if u == w or (min(u, w), max(u, w)) in edges:
                ok = False
                break
            edges.add((min(u, w), max(u, w)))
        if ok and len(edges) == n * degree // 2:
            return n, sorted(edges)
    # fallback: cycle + chords
    edges = {(i, (i + 1) % n) for i in range(n)}
    return n, sorted((min(a, b), max(a, b)) for a, b in edges)


def tseitin_cnf(n: int, edges: list, charge: dict) -> list:
    """Tseitin CNF over edge variables (1-indexed). For each vertex, clauses
    enforce parity of incident edges = charge[v]."""
    evar = {}
    for i, (u, v) in enumerate(edges):
        evar[(u, v)] = i + 1
        evar[(v, u)] = i + 1
    inc = {v: [] for v in range(n)}
    for (u, v) in edges:
        inc[u].append(evar[(u, v)])
        inc[v].append(evar[(u, v)])
    clauses = []
    for v in range(n):
        vars_ = inc[v]
        k = len(vars_)
        tgt = charge.get(v, 0)
        for mask in range(1 << k):
            if bin(mask).count("1") % 2 != tgt:
                cl = []
                for j, var in enumerate(vars_):
                    cl.append(-var if (mask >> j) & 1 else var)
                clauses.append(frozenset(cl))
    return clauses


# ---------- Resolution by Davis-Putnam variable elimination ----------

def resolve(c1: frozenset, c2: frozenset, var: int):
    """Resolvent of c1,c2 on var (c1 has +var, c2 has -var). None if tautology."""
    r = (c1 - {var}) | (c2 - {-var})
    for lit in r:
        if -lit in r:
            return None  # tautology
```

```python
    return frozenset(r)


def dp_refutation_phi(clauses: list, rng: random.Random):
    """Run DP elimination in min-occurrence order; track the active clause DB
    after each elimination. Returns (phi_count, phi_weight, n_steps, derived_empty).

    phi_count  = sum_t |DB_t|             (proofStateEntropy = card)
    phi_weight = sum_t sum_{C in DB_t}|C| (totalLiteralWeight version)
    """
    db = set(clauses)
    variables = set(abs(l) for c in db for l in c)
    phi_count = len(db)
    phi_weight = sum(len(c) for c in db)
    n_steps = 1
    derived_empty = frozenset() in db
    MAX_DB = 200000  # guard against blow-up on bad instances

    while variables and not derived_empty:
        # min-occurrence elimination order (standard good heuristic)
        occ = {v: 0 for v in variables}
        for c in db:
            for l in c:
                if abs(l) in occ:
                    occ[abs(l)] += 1
        var = min(variables, key=lambda v: occ[v])
        variables.discard(var)

        pos = [c for c in db if var in c]
        neg = [c for c in db if -var in c]
        others = [c for c in db if var not in c and -var not in c]

        resolvents = set()
        for cp in pos:
            for cn in neg:
                r = resolve(cp, cn, var)
                if r is not None:
                    resolvents.add(r)
                    if len(r) == 0:
                        derived_empty = True
        new_db = set(others) | resolvents
        # forward subsumption (keep minimal clauses) — cheap pass
        db = new_db
        phi_count += len(db)
        phi_weight += sum(len(c) for c in db)
        n_steps += 1
        if len(db) > MAX_DB:
            break

    return phi_count, phi_weight, n_steps, derived_empty


# ---------- Separator (cheap balanced-ish) ----------

def approx_separator_size(n: int, edges: list) -> int:
    """Cheap proxy for balanced separator size: min vertices whose removal
```

```python
    splits the graph ~in half. We use a simple BFS-layer cut."""
    adj = {v: set() for v in range(n)}
    for (u, v) in edges:
        adj[u].add(v); adj[v].add(u)
    # BFS from 0, cut at the layer crossing the median
    from collections import deque
    seen = {0}; layer = {0: 0}; q = deque([0])
    while q:
        x = q.popleft()
        for y in adj[x]:
            if y not in seen:
                seen.add(y); layer[y] = layer[x] + 1; q.append(y)
    if len(seen) < n:  # disconnected
        return 0
    half = n // 2
    # vertices on the boundary between the first ~half and the rest
    order = sorted(range(n), key=lambda v: layer.get(v, 0))
    side_a = set(order[:half])
    cut = set()
    for (u, v) in edges:
        if (u in side_a) != (v in side_a):
            cut.add(u if u not in side_a else v)
    return len(cut)


# ---------- Trial protocol ----------

def run_trial(seed: int) -> dict:
    rng = random.Random(seed)
    # sweep several small sizes; report the scaling exponent of Phi vs n
    ns = [6, 8, 10, 12, 14]
    data = []
    for n0 in ns:
        n, edges = random_regular_graph(n0, 3, rng)
        charge = {v: 0 for v in range(n)}
        charge[0] = 1  # odd total -> UNSAT
        clauses = tseitin_cnf(n, edges, charge)
        phi_c, phi_w, steps, empty = dp_refutation_phi(clauses, rng)
        sep = approx_separator_size(n, edges)
        data.append({"n": n, "m_edges": len(edges), "sep": sep,
                     "phi_count": phi_c, "phi_weight": phi_w,
                     "steps": steps, "derived_empty": empty})
    # scaling: log-log slope of phi_count vs n
    xs = [math.log(d["n"]) for d in data if d["phi_count"] > 0]
    ys = [math.log(d["phi_count"]) for d in data if d["phi_count"] > 0]
    if len(xs) >= 2:
        mx = sum(xs) / len(xs); my = sum(ys) / len(ys)
        num = sum((x - mx) * (y - my) for x, y in zip(xs, ys))
        den = sum((x - mx) ** 2 for x in xs)
        slope = num / den if den else 0.0
    else:
        slope = 0.0
    biggest = data[-1]
    # Sub-conjecture under test (SC1): Phi >= sep^2 on the largest instance
    # (a concrete, falsifiable scaling claim derived from the C-003b thesis
    # that cumulative entropy is forced by the separator).
```

```
    holds = biggest["phi_count"] >= max(1, biggest["sep"]) ** 2
    return {
        "metric_name": "phi_count_loglog_slope_vs_n",
        "metric_value": round(slope, 4),
        "instances_tested": len(ns),
        "n_max": max(ns),
        "conjecture_holds": holds,
        "counterexample": "" if holds else
            f"n={biggest['n']} sep={biggest['sep']} phi={biggest['phi_count']} < sep^2",
        "detail": data,
    }


if __name__ == "__main__":
    seeds = [int(x) for x in sys.argv[1:]] or [11, 23, 37]
    results = []
    for s in seeds:
        r = run_trial(s)
        results.append(r)
        print("TRIAL: " + json.dumps(r))
    slopes = [r["metric_value"] for r in results]
    holds = sum(1 for r in results if r["conjecture_holds"])
    mean = sum(slopes) / len(slopes)
    if holds == len(results):
        print(f"RESULT: SUPPORTED mean_slope={mean:.4f} support_fraction=1.0")
    elif holds == 0:
        print(f"RESULT: FALSIFIED counterexample=\"phi < sep^2\" first_failing_seed={seeds[0]}")
    else:
        print(f"RESULT: INCONCLUSIVE mixed holds={holds}/{len(results)} mean_slope={mean:.4f}")
```

6. Per-seed results

Seed	Metric value	Holds?	Counterexample
?	2.4348	□	
?	2.0177	□	
?	2.0528	□	
?	2.4751	□	
?	2.4691	□	
?	2.2173	□	
?	2.2983	□	
?	2.7629	□	
?	2.3782	□	
?	2.7571	□	
?	2.6234	□	
?	2.5729	□	
?	2.3297	□	

Aggregate statistics:

Statistic	Value
n_seeds	13
metric_mean	2.4145615384615384
metric_std	0.23518233897149377

Statistic	Value
metric_ci95_half	0.13045568957619594
metric_min	2.0177
metric_max	2.7629
support_fraction	1.0

7. Test stdout (last 2KB)


```
"m_edges": 15, "sep": 3, "phi_count": 591, "phi_weight": 2826, "steps": 16, "derived_empty": true}, {'
TRIAL: {"metric_name": "phi_count_loglog_slope_vs_n", "metric_value": 2.2983, "instances_tested": 5,
TRIAL: {"metric_name": "phi_count_loglog_slope_vs_n", "metric_value": 2.7629, "instances_tested": 5,
TRIAL: {"metric_name": "phi_count_loglog_slope_vs_n", "metric_value": 2.3782, "instances_tested": 5,
```


8. Critic adversarial review

Critic verdict: `CHALLENGE`

Critic reasoning:

> The test code focuses on the cumulative proof entropy for Tseitin formulas, which is a specific aspect of resolution proof entropy. However, the conjecture relates to the minimal rank of the associated formal power series F(φ), not just the entropy. The test does not provide evidence about the rank of the formal power series or its relation to resolution proof entropy.

9. Final verdict & safety rail

Verdict: `INCONCLUSIVE`

Reasoning:

> The test focuses on cumulative proof entropy for Tseitin formulas, which is a specific aspect of resolution proof entropy. However, the conjecture concerns the minimal rank of the associated formal power series F(φ), not just the entropy. The test does not provide evidence about the rank of the formal power series or its relation to resolution proof entropy. | next: Further investigation is needed to validate the conjecture by examining the relationship between the minimal rank of the formal pow

11. Audit log (LLM calls)

Total LLM calls: 6

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
1	propose	ollama_remote	glm4:latest	0	0	20296
2	preregistration	ollama_remote	glm4:latest	0	0	15511
3	novelty	ollama_remote	glm4:latest	0	0	8823
4	novelty	ollama_remote	glm4:latest	0	0	16778
5	critic	ollama_remote	glm4:latest	0	0	11750
6	judge	ollama_remote	glm4:latest	0	0	9749

Totals: 0 input tokens, 0 output tokens, ~$0.0000 cost, 82907 ms total latency. Provider mix: {'ollama_remote': 6}

(full prompt+response transcripts available in `research/audit/988964c281c1.jsonl`)

12. Reproducibility

To reproduce this cycle:

1. Apply the test harness in section 5.1 with seeds 11,23,37,53,71
2. The Python sandbox is pure-stdlib; any Python 3.11+ should suffice
3. The full LLM transcripts are in `research/audit/988964c281c1.jsonl` for verification of provenance
4. The pre-registration hash in section 2 commits to the test design before execution; tampering would be detectable.

Sandbox file: `research/pvsnp_sandbox/test_*.py` (per-cycle)
Replay tarball: `research/replay/988964c281c1.tar.gz` (if generated)