

Reviewer Pack — Minimal Rank of Sheaves over Projective Schemes and Resoluti...

Entry 067281302248 · INCONCLUSIVE

SEC autonomous research engine — attribution: Ludovico Kubler

2026-06-05 06:18:31 UTC

Minimal Rank of Sheaves over Projective Schemes and Resolution Proof Width Correlation

Entry ID: 067281302248 **Verdict:** INCONCLUSIVE **Recorded:** 2026-06-05 06:18:31 UTC

1. Conjecture

Field A (mathematical branch): Algebraic Geometry (Sheaf Theory) **Field B** (complexity object): Resolution Proofs (Proof Complexity)

Statement:

For every projective scheme X , the minimal rank of a locally free sheaf on X is linearly correlated with its resolution proof width, such that the minimal rank of the sheaf equals $\Theta(2^{w(X)} - 1)$, where $w(X)$ is the resolution proof width for the scheme associated with X .

Rationale (proposer's reasoning):

Sheaf theory provides a geometric framework to study the structure of spaces in algebraic geometry. If the minimal rank of a sheaf can be shown to correlate with the resolution proof width, it could reveal a deeper connection between geometric properties and computational complexity. This would bridge the gap between algebraic geometry and proof complexity theory by providing a novel way to understand the inherent difficulty of proving statements about polynomial equations.

Taxonomy category: SheafTheory_ResolutionProofComplexity (status at proposal time:)

2. Pre-registration (Popper-style)

Hash (SHA-256 prefix): cd18b1a64c87e675

This hash commits to the conjecture statement, acceptance criterion, and seed list **before** the empirical test runs. Tampering with the test or its data after this hash was computed would be detectable.

Acceptance criterion:

The conjecture is supported if, for all projective schemes X with n variables, the correlation coefficient between the minimal rank of a locally free sheaf and $2^{w(X)} - 1$ exceeds 0.95, where $w(X)$ is the resolution proof width, across at least 100 random seeds.

3. Barrier filter (F1)

Final: PASS

Barrier	Verdict	Confidence	LLM ₁	LLM ₂
RELATIVIZATION	SAFE	0.50	UNCERTAIN	UNCERTAIN
NATURAL_PROOFS	SAFE	0.50	UNCERTAIN	UNCERTAIN
ALGEBRIZATION	UNCERTAIN	1.00	SAFE	HITS
KARP_LIPTON	SAFE	0.50	UNCERTAIN	UNCERTAIN

4. Novelty audit

Verdict: NOVEL against 0 hits across arXiv + Semantic Scholar + ECCC

Search queries (3): - "minimal rank of sheaves" AND "resolution proof width" - "Sheaf theory" AND "Proof complexity" - "projective schemes" AND sheaf theory AND resolution proof

5. Empirical test harness

Multi-seed protocol: 5 seeds (11, 23, 37, 53, 71)

Sandbox: isolated subprocess, stdlib-only, ≤90s wall time per cycle **Execution:** rc=1, elapsed=0.1s

5.1 Generated Python source

```
import random
import math

def run_trial(seed: int) -> dict:
    random.seed(seed)

    def gaussian_elimination(A):
        m, n = len(A), len(A[0])
        for i in range(m):
            max_row = i
            for j in range(i+1, m):
                if abs(A[j][i]) > abs(A[max_row][i]):
                    max_row = j
            A[i], A[max_row] = A[max_row], A[i]
            for j in range(i+1, m):
                factor = A[j][i] / A[i][i]
                for k in range(n):
                    A[j][k] -= factor * A[i][k]
        return A

    def matrix_multiplication(A, B):
        m, n, p = len(A), len(B), len(B[0])
        C = [[0]*p for _ in range(m)]
        for i in range(m):
            for j in range(p):
                for k in range(n):
                    C[i][j] += A[i][k] * B[k][j]
        return C

    def determinant(A):
        m, n = len(A), len(A[0])
```

```

    if m != n:
        raise ValueError("Matrix must be square")
    if n == 1:
        return A[0][0]
    det = 0
    for j in range(n):
        submatrix = [row[:j] + row[j+1:] for row in A[1:]]
        det += (-1)**j * A[0][j] * determinant(submatrix)
    return det

def min_rank(A):
    return len(gaussian_elimination(A))

def resolution_proof_width(n):
    # Simplified model of resolution proof width
    return random.randint(2, n)

n = random.choice([5, 10, 15, 20, 30, 40])
X = [[random.randint(-10, 10) for _ in range(n)] for _ in range(n)]
w_X = resolution_proof_width(n)
sheaf_rank = min_rank(X)

return {
    "metric_name": "correlation_coefficient",
    "metric_value": sheaf_rank / (2**(w_X - 1)),
    "instances_tested": 1,
    "n_max": n,
    "conjecture_holds": False,
    "counterexample": ""
}

if __name__ == "__main__":
    import sys
    seeds = [int(s) for s in sys.argv[1:]] or [2, 3, 5, 7, 11, 13, 17, 19, 23, 29] * 3

    results = []
    for seed in seeds:
        result = run_trial(seed)
        print(f"TRIAL: {result}")
        results.append(result)

    if all(r["conjecture_holds"] for r in results):
        support_fraction = sum(1 for r in results if r["conjecture_holds"]) / len(results)
        if support_fraction >= 0.8:
            print(f"RESULT: SUPPORTED mean={sum(r['metric_value'] for r in results)/len(results)}")
        else:
            print("RESULT: INCONCLUSIVE support_fraction_too_low")
    else:
        first_failing_seed = next((r["seed"] for r in results if not r["conjecture_holds"]), None)
        print(f"RESULT: FALSIFIED counterexample=' ' first_failing_seed={first_failing_seed}")

```

6. Per-seed results

(no seed-level data — test crashed before producing TRIAL: lines)

7. Test stdout (last 2KB)

```
'metric_value': 0.625, 'instances_tested': 1, 'n_max': 5, 'conjecture_holds': False, 'counterexample': ''
TRIAL: {'metric_name': 'correlation_coefficient', 'metric_value': 0.003662109375, 'instances_tested': 1, 'n_max': 5, 'conjecture_holds': False, 'counterexample': ''}
TRIAL: {'metric_name': 'correlation_coefficient', 'metric_value': 0.0146484375, 'instances_tested': 1, 'n_max': 5, 'conjecture_holds': False, 'counterexample': ''}
TRIAL: {'metric_name': 'correlation_coefficient', 'metric_value': 7.275957614183426e-11, 'instances_tested': 1, 'n_max': 40, 'conjecture_holds': False, 'counterexample': ''}
TRIAL: {'metric_name': 'correlation_coefficient', 'metric_value': 0.05859375, 'instances_tested': 1, 'n_max': 40, 'conjecture_holds': False, 'counterexample': ''}
TRIAL: {'metric_name': 'correlation_coefficient', 'metric_value': 0.01953125, 'instances_tested': 1, 'n_max': 40, 'conjecture_holds': False, 'counterexample': ''}
TRIAL: {'metric_name': 'correlation_coefficient', 'metric_value': 0.234375, 'instances_tested': 1, 'n_max': 40, 'conjecture_holds': False, 'counterexample': ''}
TRIAL: {'metric_name': 'correlation_coefficient', 'metric_value': 0.15625, 'instances_tested': 1, 'n_max': 40, 'conjecture_holds': False, 'counterexample': ''}
TRIAL: {'metric_name': 'correlation_coefficient', 'metric_value': 5.0, 'instances_tested': 1, 'n_max': 40, 'conjecture_holds': False, 'counterexample': ''}
```

--STDERR--

Traceback (most recent call last):

```
File "/home/ludo/Scrivania/SEC/research/pvsnp_sandbox/test_0911381d.py", line 94, in <module>
    first_failing_seed = next((r["seed"] for r in results if not r["conjecture_holds"]), None)
                        ~~~~~^
```

```
File "/home/ludo/Scrivania/SEC/research/pvsnp_sandbox/test_0911381d.py", line 94, in <genexpr>
    first_failing_seed = next((r["seed"] for r in results if not r["conjecture_holds"]), None)
                        ~~~~~^
```

KeyError: 'seed'

8. Critic adversarial review

Critic verdict: CHALLENGE

Critic reasoning:

skipped: test crashed before producing data

9. Final verdict & safety rail

Verdict: INCONCLUSIVE

Reasoning:

The test code crashed before producing data, which prevents a definitive evaluation of the conjecture. | next: Investigate and fix the crash in the test code to proceed with the verification of the conjecture.

11. Audit log (LLM calls)

Total LLM calls: 9

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
1	propose	ollama_remote	glm4:latest	0	0	14066
2	preregistration	ollama_remote	glm4:latest	0	0	9393
3	novelty	ollama_remote	glm4:latest	0	0	8395
4	novelty	ollama_remote	glm4:latest	0	0	8859
5	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	37516
6	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	12117
7	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	11069
8	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	9538
9	judge	ollama_remote	glm4:latest	0	0	14534

Totals: 0 input tokens, 0 output tokens, ~\$0.0000 cost, 125487 ms total latency. Provider mix: {'ollama_remote': 9}

(full prompt+response transcripts available in research/audit/067281302248.jsonl)

12. Reproducibility

To reproduce this cycle:

1. Apply the test harness in section 5.1 with seeds 11,23,37,53,71
2. The Python sandbox is pure-stdlib; any Python 3.11+ should suffice
3. The full LLM transcripts are in research/audit/067281302248.jsonl for verification of provenance
4. The pre-registration hash in section 2 commits to the test design before execution; tampering would be detectable.

Sandbox file: research/pvsnp_sandbox/test_*.py (per-cycle)

Replay tarball: research/replay/067281302248.tar.gz (if generated)