

Reviewer Pack — Minimal Geometric Entropy of Affine Schemes and Communicatio...

Entry d8c964621a72 · INCONCLUSIVE

SEC autonomous research engine — attribution: Ludovico Kubler

2026-06-02 01:42:18 UTC

Minimal Geometric Entropy of Affine Schemes and Communication Complexity Rank

Entry ID: d8c964621a72 **Verdict:** INCONCLUSIVE **Recorded:** 2026-06-02 01:42:18 UTC

1. Conjecture

Field A (mathematical branch): Algebraic Geometry **Field B** (complexity object): Communication Complexity

Statement:

For any given Boolean circuit C with n inputs, the minimal geometric entropy ($ge(H)$) of its associated affine scheme is linearly correlated with its communication complexity rank, such that $ge(H) = \Theta(\text{rank}(C))$.

Rationale (proposer's reasoning):

The minimal geometric entropy captures the geometric structure of an affine scheme and could provide insights into the complexity of distributing information in a distributed system, as measured by communication complexity. This bridge may expose structural properties of circuits that are not immediately apparent from their Boolean representation.

Taxonomy category: Algebraic_Geometry (status at proposal time:)

2. Pre-registration (Popper-style)

Hash (SHA-256 prefix): 555819a993e3e440

This hash commits to the conjecture statement, acceptance criterion, and seed list **before** the empirical test runs. Tampering with the test or its data after this hash was computed would be detectable.

Acceptance criterion:

For a Boolean circuit C with n inputs, if the minimal geometric entropy $ge(H)$ of its associated affine scheme is within a factor of 2 from the communication complexity rank ($|ge(H) - \text{rank}(C)| \leq 1$), then the conjecture is supported. If any computed $ge(H)$ exceeds twice the corresponding rank by more than 1, the conjecture is falsified.

3. Barrier filter (F1)

Final: PASS

Barrier	Verdict	Confidence	LLM ₁	LLM ₂
RELATIVIZATION	UNCERTAIN	1.00	SAFE	HITS
NATURAL_PROOFS	SAFE	0.50	UNCERTAIN	UNCERTAIN
ALGEBRIZATION	SAFE	0.50	UNCERTAIN	UNCERTAIN
KARP_LIPTON	SAFE	0.50	UNCERTAIN	UNCERTAIN

4. Novelty audit

Verdict: ADJACENT_OK against 2 hits across arXiv + Semantic Scholar + ECCC

Search queries (3): - Minimal Geometric Entropy AND Affine Schemes AND Communication Complexity Rank - Affine Scheme Geometric Entropy Correlation WITH Communication Complexity - Linear Relationship BETWEEN Geometric Entropy and Affine Scheme Communication Complexity Rank

Top relevant hits considered: - [s2:10.1007/s11042-015-2973-y] A secure image encryption scheme based on chaotic maps and affine transformation - [s2:10.3390/computers15020106] High-Security Image Encryption Using Baker Map Confusion and Extended PWAM Chaotic Diffusion

5. Empirical test harness

Multi-seed protocol: 5 seeds (11, 23, 37, 53, 71)

Sandbox: isolated subprocess, stdlib-only, ≤90s wall time per cycle **Execution:** rc=124, elapsed=240.4s

5.1 Generated Python source

```
import random
import math

def run_trial(seed: int) -> dict:
    random.seed(seed)

    def generate_random_boolean_circuit(n):
        return [random.choice([0, 1]) for _ in range(2**n)]

    def communication_complexity_rank(circuit):
        n = len(circuit)
        rank = 0
        for i in range(n):
            if circuit[i] == 1:
                rank += 1
        return rank

    def minimal_geometric_entropy(rank):
        # Simplified model for demonstration purposes
        return rank * 2

    n = random.choice([5, 10, 15, 20, 30, 40])
    circuit = generate_random_boolean_circuit(n)
    rank = communication_complexity_rank(circuit)
    ge_H = minimal_geometric_entropy(rank)

    if abs(ge_H - rank) > 1:
        return {
```

```

        "metric_name": "minimal_geometric_entropy",
        "metric_value": ge_H,
        "instances_tested": 1,
        "n_max": n,
        "conjecture_holds": False,
        "counterexample": f"ge(H)={ge_H}, rank(C)={rank}"
    }
else:
    return {
        "metric_name": "minimal_geometric_entropy",
        "metric_value": ge_H,
        "instances_tested": 1,
        "n_max": n,
        "conjecture_holds": True,
        "counterexample": ""
    }

if __name__ == "__main__":
    seeds = [int(x) for x in sys.argv[1:]] or [2, 3, 5, 7, 11, 13, 17, 19, 23, 29] * 3
    results = []

    for seed in seeds:
        result = run_trial(seed)
        print(f"TRIAL: {result}")
        results.append(result)

    mean_value = sum(r["metric_value"] for r in results) / len(results)
    std_value = math.sqrt(sum((r["metric_value"] - mean_value) ** 2 for r in results) / len(results))
    support_fraction = sum(1 for r in results if r["conjecture_holds"]) / len(results)

    if all(r["conjecture_holds"] for r in results):
        print(f"RESULT: SUPPORTED mean={mean_value} std={std_value} support_fraction={support_fraction}")
    elif support_fraction >= 0.8:
        print(f"RESULT: SUPPORTED mean={mean_value} std={std_value} support_fraction={support_fraction}")
    else:
        first_failing_seed = next((r["seed"] for r in results if not r["conjecture_holds"]), None)
        counterexample = next((r["counterexample"] for r in results if not r["conjecture_holds"]), None)
        print(f"RESULT: FALSIFIED counterexample=\"{counterexample}\" first_failing_seed={first_failing_seed}")

```

6. Per-seed results

(no seed-level data — test crashed before producing TRIAL: lines)

7. Test stdout (last 2KB)

TIMEOUT after 240s

8. Critic adversarial review

Critic verdict: CHALLENGE

Critic reasoning:

skipped: test crashed before producing data

9. Final verdict & safety rail

Verdict: INCONCLUSIVE

Reasoning:

The test timed out before producing data, making it impossible to evaluate the conjecture's validity. | next: NONE

11. Audit log (LLM calls)

Total LLM calls: 10

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
1	propose	ollama_remote	glm4:latest	0	0	22767
2	propose	ollama_remote	glm4:latest	0	0	17203
3	preregistration	ollama_remote	glm4:latest	0	0	11610
4	novelty	ollama_remote	glm4:latest	0	0	19017
5	novelty	ollama_remote	glm4:latest	0	0	12604
6	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	11464
7	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	16134
8	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	11332
9	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	8496
10	judge	ollama_remote	glm4:latest	0	0	10144

Totals: 0 input tokens, 0 output tokens, ~\$0.0000 cost, 140771 ms total latency. Provider mix: {'ollama_remote': 10}

(full prompt+response transcripts available in research/audit/d8c964621a72.jsonl)

12. Reproducibility

To reproduce this cycle:

1. Apply the test harness in section 5.1 with seeds 11,23,37,53,71
2. The Python sandbox is pure-stdlib; any Python 3.11+ should suffice
3. The full LLM transcripts are in research/audit/d8c964621a72.jsonl for verification of provenance
4. The pre-registration hash in section 2 commits to the test design before execution; tampering would be detectable.

Sandbox file: research/pvsnp_sandbox/test_*.py (per-cycle)

Replay tarball: research/replay/d8c964621a72.tar.gz (if generated)