

Reviewer Pack — Minimal Rank of Geometric Langlands Correspondence over SAT ...

Entry e0b483106b38 · INCONCLUSIVE

SEC autonomous research engine — attribution: Ludovico Kubler

2026-05-27 11:50:04 UTC

Minimal Rank of Geometric Langlands Correspondence over SAT Formulas

Entry ID: e0b483106b38 **Verdict:** INCONCLUSIVE **Recorded:** 2026-05-27 11:50:04 UTC

1. Conjecture

Field A (mathematical branch): Geometric Langlands Theory **Field B** (complexity object): Complexity Theory: Boolean Circuit Satisfiability

Statement:

[‘For every k -CNF formula F with n variables, the minimal rank of the geometric Langlands lattice associated with F is $\Theta(n^{1/4} \log n)$.’, ‘Equivalently, for any polynomial-time computable function $f(n)$, there exists a k -CNF formula F such that the minimal rank of its geometric Langlands lattice is at least $n^{1/4} \log n$ but not higher than $f(n)$.’]

Rationale (proposer’s reasoning):

[‘The Geometric Langlands correspondence relates algebraic varieties to certain automorphic forms, which are objects in representation theory. If this conjecture holds, it would imply a direct connection between the geometric complexity of an object and its representation-theoretic properties, providing new insights into both fields.’, ‘This might expose hidden structures in SAT formulas that could lead to improved algorithms for solving them, as automorphic forms often have special properties that can be exploited.’]

Taxonomy category: Geometric_Langlands_Theory (status at proposal time:)

2. Pre-registration (Popper-style)

Hash (SHA-256 prefix): 431d646773bf2b53

This hash commits to the conjecture statement, acceptance criterion, and seed list **before** the empirical test runs. Tampering with the test or its data after this hash was computed would be detectable.

Acceptance criterion:

The conjecture is supported if at least 80% of the generated k -CNF formulas with n variables ($n \leq 40$) have a geometric Langlands lattice minimal rank within 10% of $n^{1/4} \log n$, calculated across 30 seeds. The conjecture is falsified if any seed produces a geometric Langlands lattice minimal rank outside this range.

3. Barrier filter (F1)

Final: PASS

Barrier	Verdict	Confidence	LLM ₁	LLM ₂
RELATIVIZATION	SAFE	0.50	UNCERTAIN	UNCERTAIN
NATURAL_PROOFS	SAFE	0.50	UNCERTAIN	UNCERTAIN
ALGEBRIZATION	SAFE	0.50	UNCERTAIN	UNCERTAIN
KARP_LIPTON	SAFE	0.50	UNCERTAIN	UNCERTAIN

4. Novelty audit

Verdict: NOVEL against 0 hits across arXiv + Semantic Scholar + ECCC

Search queries (3): - geometric langlands theory AND satisfiability modulo theories - CNF formulas AND geometric langlands lattices AND complexity - minimal rank AND geometric langlands AND circuit satisfiability

5. Empirical test harness

Multi-seed protocol: 5 seeds (11, 23, 37, 53, 71)

Sandbox: isolated subprocess, stdlib-only, ≤90s wall time per cycle **Execution:** rc=0, elapsed=0.0s

5.1 Generated Python source

```
import random
import math

def run_trial(seed: int) -> dict:
    random.seed(seed)

    def generate_kcnf(n, k):
        clauses = []
        for _ in range(k):
            variables = random.sample(range(1, n+1), 3)
            clause = [random.choice([var, -var]) for var in variables]
            clauses.append(clause)
        return clauses

    def geometric_langlands_lattice(clauses):
        # Construct a lattice based on the clause structure
        # This is a placeholder implementation
        lattice = []
        for clause in clauses:
            for literal in clause:
                if literal not in lattice:
                    lattice.append(literal)
        return lattice

    def min_rank(lattice):
        # Placeholder function to estimate the minimal rank of the lattice
        # For simplicity, we use the number of unique literals as a proxy
        return len(set(abs(x) for x in lattice))

    n = random.randint(5, 40)
```

```

k = max(1, min(n // 2, 3)) # Ensure at least one clause and not too many
formula = generate_kcnf(n, k)
lattice = geometric_langlands_lattice(formula)
rank = min_rank(lattice)

expected_rank = n ** (1/4) * math.log(n)
tolerance = 0.1 * expected_rank

metric_name = "Minimal Rank"
metric_value = rank
instances_tested = 1
conjecture_holds = abs(rank - expected_rank) <= tolerance
counterexample = "" if conjecture_holds else f"Rank {rank} does not match expected {expected_rank}"

return {
    "metric_name": metric_name,
    "metric_value": metric_value,
    "instances_tested": instances_tested,
    "conjecture_holds": conjecture_holds,
    "counterexample": counterexample
}

if __name__ == "__main__":
    seeds = [int(arg) for arg in sys.argv[1:]] or [2, 3, 5, 7, 11, 13, 17, 19, 23, 29, 31, 37, 41, 43, 47, 53, 59, 61, 67, 71, 73, 79, 83, 89, 97]

    results = []
    for seed in seeds:
        trial_result = run_trial(seed)
        print(f"TRIAL: {trial_result}")
        results.append(trial_result)

    mean_rank = sum(result["metric_value"] for result in results) / len(results)
    std_rank = math.sqrt(sum((result["metric_value"] - mean_rank) ** 2 for result in results) / len(results))
    support_fraction = sum(1 for result in results if result["conjecture_holds"]) / len(results)

    if support_fraction >= 0.8:
        print(f"RESULT: SUPPORTED mean={mean_rank:.4f} std={std_rank:.4f} support_fraction={support_fraction:.4f}")
    elif any(not result["conjecture_holds"] for result in results):
        first_failing_seed = next(seed for seed, result in zip(seeds, results) if not result["conjecture_holds"])
        print(f"RESULT: FALSIFIED counterexample=\"Rank out of tolerance\" first_failing_seed={first_failing_seed}")
    else:
        print("RESULT: INCONCLUSIVE insufficient_data")

```

6. Per-seed results

(no seed-level data — test crashed before producing TRIAL: lines)

7. Test stdout (last 2KB)

```

'Minimal Rank', 'metric_value': 8, 'instances_tested': 1, 'conjecture_holds': False, 'counterexample':
TRIAL: {'metric_name': 'Minimal Rank', 'metric_value': 7, 'instances_tested': 1, 'conjecture_holds':
TRIAL: {'metric_name': 'Minimal Rank', 'metric_value': 7, 'instances_tested': 1, 'conjecture_holds':
TRIAL: {'metric_name': 'Minimal Rank', 'metric_value': 6, 'instances_tested': 1, 'conjecture_holds':
TRIAL: {'metric_name': 'Minimal Rank', 'metric_value': 8, 'instances_tested': 1, 'conjecture_holds':
TRIAL: {'metric_name': 'Minimal Rank', 'metric_value': 8, 'instances_tested': 1, 'conjecture_holds':
TRIAL: {'metric_name': 'Minimal Rank', 'metric_value': 6, 'instances_tested': 1, 'conjecture_holds':

```

```

TRIAL: {'metric_name': 'Minimal Rank', 'metric_value': 7, 'instances_tested': 1, 'conjecture_holds':
TRIAL: {'metric_name': 'Minimal Rank', 'metric_value': 7, 'instances_tested': 1, 'conjecture_holds':
TRIAL: {'metric_name': 'Minimal Rank', 'metric_value': 8, 'instances_tested': 1, 'conjecture_holds':
TRIAL: {'metric_name': 'Minimal Rank', 'metric_value': 8, 'instances_tested': 1, 'conjecture_holds':
TRIAL: {'metric_name': 'Minimal Rank', 'metric_value': 8, 'instances_tested': 1, 'conjecture_holds':
RESULT: FALSIFIED counterexample="Rank out of tolerance" first_failing_seed=11

```

8. Critic adversarial review

Critic verdict: CHALLENGE

Critic reasoning:

The empirical test has only tested a small number of instances ($n \leq 15$). The metric does not scale trivially with n , but the limited sample size may not be representative. Additionally, there is no evidence that the conjecture holds for larger values of n , which could indicate a failure mode such as metric saturation or selection bias.

9. Final verdict & safety rail

Verdict: INCONCLUSIVE

Reasoning:

Safety rail: critic_challenge_falsified | original: The test results show that the geometric Langlands lattice minimal rank does not match the expected value within the tolerance range for at least one | next: Further investigation is needed to understand why the conjecture fails. This may involve analyzing the specific properties of the counterexamples or exploring alternative models that could better capture the behavior of the geometric Langlands lattice.

11. Audit log (LLM calls)

Total LLM calls: 11

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
1	propose	ollama_remote	glm4:latest	0	0	11450
2	propose	ollama_remote	glm4:latest	0	0	10573
3	preregistration	ollama_remote	glm4:latest	0	0	6244
4	novelty	ollama_remote	glm4:latest	0	0	4532
5	novelty	ollama_remote	glm4:latest	0	0	5586
6	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	11553
7	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	8144
8	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	6907
9	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	9563
10	critic	ollama_remote	glm4:latest	0	0	9316
11	judge	ollama_remote	glm4:latest	0	0	8949

Totals: 0 input tokens, 0 output tokens, ~\$0.0000 cost, 92818 ms total latency. Provider mix: {'ollama_remote': 11}

(full prompt+response transcripts available in `research/audit/e0b483106b38.jsonl`)

12. Reproducibility

To reproduce this cycle:

1. Apply the test harness in section 5.1 with seeds 11,23,37,53,71
2. The Python sandbox is pure-stdlib; any Python 3.11+ should suffice
3. The full LLM transcripts are in `research/audit/e0b483106b38.jsonl` for verification of provenance
4. The pre-registration hash in section 2 commits to the test design before execution; tampering would be detectable.

Sandbox file: `research/pvsnp_sandbox/test_*.py` (per-cycle)

Replay tarball: `research/replay/e0b483106b38.tar.gz` (if generated)