

Reviewer Pack — Minimal Number of Toric Resolutions and Frege Proof Length

Entry f5a952521849 · INCONCLUSIVE

SEC autonomous research engine — attribution: Ludovico Kubler

2026-06-03 23:28:04 UTC

Minimal Number of Toric Resolutions and Frege Proof Length

Entry ID: f5a952521849 **Verdict:** INCONCLUSIVE **Recorded:** 2026-06-03 23:28:04 UTC

1. Conjecture

Field A (mathematical branch): Toric Geometry **Field B** (complexity object): Frege Proof Complexity

Statement:

For any given Boolean formula φ with n variables, the length of the shortest Frege proof for φ is linearly correlated with the minimum number of toric resolutions required to construct the associated algebraic variety, such that $|A(\varphi)| = \Theta(\text{num_resolutions}(\varphi))$, where $\text{num_resolutions}(\varphi)$ is a function measuring the minimum number of irreducible components in the toric resolution of the variety corresponding to φ .

Rationale (proposer's reasoning):

Toric geometry provides a bridge between algebraic geometry and Boolean satisfiability problems. The structure of an algebraic variety can potentially reflect the complexity of proving its satisfiability, as Frege proofs correspond to paths in the variety's realization space.

Taxonomy category: ToricGeometry x FregeProofComplexity (status at proposal time:)

2. Pre-registration (Popper-style)

Hash (SHA-256 prefix): 7d8922a81fd6c643

This hash commits to the conjecture statement, acceptance criterion, and seed list **before** the empirical test runs. Tampering with the test or its data after this hash was computed would be detectable.

Acceptance criterion:

The conjecture is supported if, for at least 24 out of 30 random seeds, the correlation coefficient between the number of minimal toric resolutions and the length of the shortest Frege proof for each Boolean formula φ is $|r| \geq 0.7$, where r is calculated using Pearson's correlation coefficient.

3. Barrier filter (F1)

Final: PASS

Barrier	Verdict	Confidence	LLM ₁	LLM ₂
RELATIVIZATION	SAFE	0.50	UNCERTAIN	UNCERTAIN
NATURAL_PROOFS	SAFE	0.50	UNCERTAIN	UNCERTAIN
ALGEBRIZATION	UNCERTAIN	1.00	UNCERTAIN	HITS
KARP_LIPTON	SAFE	0.50	UNCERTAIN	UNCERTAIN

4. Novelty audit

Verdict: NOVEL against 0 hits across arXiv + Semantic Scholar + ECCC

5. Empirical test harness

Multi-seed protocol: 5 seeds (11, 23, 37, 53, 71)

Sandbox: isolated subprocess, stdlib-only, ≤90s wall time per cycle **Execution:** rc=124, elapsed=240.4s

5.1 Generated Python source

```
import random
import math
from fractions import Fraction

def run_trial(seed: int) -> dict:
    random.seed(seed)

    def generate_boolean_formula(n):
        return ''.join(random.choice('01') for _ in range(2**n - 1))

    def count_irreducible_components(formula):
        # Placeholder function to simulate counting irreducible components
        # This is a dummy implementation and should be replaced with actual logic
        return random.randint(1, n)

    def shortest_frege_proof_length(formula):
        # Placeholder function to simulate Frege proof length calculation
        # This is a dummy implementation and should be replaced with actual logic
        return len(formula) * 2

    results = []
    for _ in range(30):
        n = random.choice([5, 10, 15, 20, 30, 40])
        formula = generate_boolean_formula(n)
        num_resolutions = count_irreducible_components(formula)
        proof_length = shortest_frege_proof_length(formula)
        results.append((num_resolutions, proof_length))

    if not results:
        return {
            "metric_name": "correlation_coefficient",
            "metric_value": 0,
            "instances_tested": 0,
            "n_max": n,
            "conjecture_holds": False,
            "counterexample": "no_results"
        }
```

```

num_resolutions = [r[0] for r in results]
proof_lengths = [r[1] for r in results]
mean_num_resolutions = sum(num_resolutions) / len(num_resolutions)
mean_proof_lengths = sum(proof_lengths) / len(proof_lengths)

numerator = sum((num_resolutions[i] - mean_num_resolutions) * (proof_lengths[i] - mean_proof_lengths) for i in range(len(results)))
denominator = math.sqrt(sum((num_resolutions[i] - mean_num_resolutions)**2 for i in range(len(results))))

if denominator == 0:
    return {
        "metric_name": "correlation_coefficient",
        "metric_value": 0,
        "instances_tested": len(results),
        "n_max": n,
        "conjecture_holds": False,
        "counterexample": "division_by_zero"
    }

correlation_coefficient = numerator / denominator

return {
    "metric_name": "correlation_coefficient",
    "metric_value": abs(correlation_coefficient),
    "instances_tested": len(results),
    "n_max": n,
    "conjecture_holds": abs(correlation_coefficient) >= 0.7,
    "counterexample": ""
}

if __name__ == "__main__":
    import sys
    if not sys.argv[1:]:
        seeds = [2, 3, 5, 7, 11, 13, 17, 19, 23, 29, 31, 37, 41, 43, 47, 53, 59, 61, 67, 71, 73, 79]
    else:
        seeds = [int(s) for s in sys.argv[1:]]

    results = []
    for seed in seeds:
        result = run_trial(seed)
        print(f"TRIAL: {result}")
        results.append(result)

    if not results:
        print("RESULT: INCONCLUSIVE no_results")
    else:
        mean_metric_value = sum(r["metric_value"] for r in results) / len(results)
        std_metric_value = math.sqrt(sum((r["metric_value"] - mean_metric_value)**2 for r in results) / len(results))
        support_fraction = sum(1 for r in results if r["conjecture_holds"]) / len(results)

        if support_fraction >= 0.8:
            print(f"RESULT: SUPPORTED mean={mean_metric_value} std={std_metric_value} support_fraction={support_fraction}")
        elif any(not r["conjecture_holds"] for r in results):
            first_failing_seed = next(seed for seed, result in zip(seeds, results) if not result["conjecture_holds"])
            print(f"RESULT: FALSIFIED counterexample=\"correlation_coefficient < 0.7\" first_failing_seed={first_failing_seed}")
        else:
            print("RESULT: UNKNOWN")

```

```
print("RESULT: INCONCLUSIVE no_support")
```

6. Per-seed results

(no seed-level data — test crashed before producing TRIAL: lines)

7. Test stdout (last 2KB)

TIMEOUT after 240s

8. Critic adversarial review

Critic verdict: CHALLENGE

Critic reasoning:

skipped: test crashed before producing data

9. Final verdict & safety rail

Verdict: INCONCLUSIVE

Reasoning:

The test timed out before producing data, which prevents us from evaluating the correlation coefficient required to support or falsify the conjecture. | next: NONE

11. Audit log (LLM calls)

Total LLM calls: 9

#	Phase	Provider	Model	Tokens in	out	Latency (ms)
1	propose	ollama_remote	glm4:latest	0	0	17400
2	preregistration	ollama_remote	glm4:latest	0	0	12777
3	novelty	ollama_remote	glm4:latest	0	0	8134
4	novelty	ollama_remote	glm4:latest	0	0	8658
5	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	10729
6	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	11228
7	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	37135
8	test_gen	ollama_remote	qwen2.5-coder:7b	0	0	45446
9	judge	ollama_remote	glm4:latest	0	0	31588

Totals: 0 input tokens, 0 output tokens, ~\$0.0000 cost, 183095 ms total latency. Provider mix: {'ollama_remote': 9}

(full prompt+response transcripts available in `research/audit/f5a952521849.jsonl`)

12. Reproducibility

To reproduce this cycle:

1. Apply the test harness in section 5.1 with seeds 11,23,37,53,71
2. The Python sandbox is pure-stdlib; any Python 3.11+ should suffice
3. The full LLM transcripts are in `research/audit/f5a952521849.jsonl` for verification of provenance
4. The pre-registration hash in section 2 commits to the test design before execution; tampering would be detectable.

Sandbox file: research/pvsnp_sandbox/test_*.py (per-cycle)

Replay tarball: research/replay/f5a952521849.tar.gz (if generated)